



UNIVERSITAT OBERTA DE CATALUNYA (UOC)

MASTER'S DEGREE IN DATA SCIENCE

MASTER'S THESIS

AREA: NATURAL LANGUAGE PROCESSING

Factible: A Multi-Agent System for Automated Fact-Checking of YouTube Videos

Author: Begoña Echavarren Sánchez

Tutor: Josep-Anton Mir Tutusaus

December 2025

Copyright

This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.



Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0)

<https://creativecommons.org/licenses/by-nc-nd/4.0/>

Source Code:

The complete source code for this project is publicly available at:

<https://github.com/begoechavarren/factable>

FINAL PROJECT RECORD

Title of the project:	Factible: A Multi-Agent System for Automated Fact-Checking of YouTube Videos
Author's name:	Begoña Echavarren Sánchez
Tutor's name:	Josep-Anton Mir Tutusaus
Delivery date:	12/2025
Degree or program:	Master's Degree in Data Science
Final Project area:	Natural Language Processing
Language of the project:	English
Keywords:	fact-checking, LLM, multi-agent systems

Acknowledgements

I would like to express my sincere gratitude to my tutor, Josep-Anton Mir Tutsaus, for his guidance, insightful feedback, and continuous support throughout this project. His expertise in data science and machine learning was invaluable in shaping this work.

I am also grateful to the Universitat Oberta de Catalunya (UOC) for providing an excellent learning environment and the resources necessary to complete this Master's Degree in Data Science.

Special thanks to the open-source community, particularly the developers of Pydantic AI, FastAPI, and the various tools that made this implementation possible. The availability of high-quality open-source software was essential to this project's success.

Finally, I would like to thank my family and friends for their patience, encouragement, and understanding during the intensive months of this thesis work.

Abstract

Misinformation on video platforms represents a growing challenge in the digital age. YouTube, with over 2 billion monthly users and 500 hours of video uploaded every minute, serves as a primary information source for millions worldwide. However, the absence of effective verification mechanisms allows false or misleading claims to spread at massive scale, impacting public health, democratic processes, and social cohesion.

This thesis presents Factible, a multi-agent system for automated fact-checking of YouTube videos. The system implements an end-to-end pipeline that processes video content through five specialized components: transcript extraction, claim detection using thesis-first reasoning, query generation for evidence retrieval, open-web evidence collection with source reliability assessment, and verdict synthesis with confidence levels.

The implementation leverages large language models (LLMs) for reasoning tasks while employing classical algorithms for deterministic operations. Key innovations include a thesis-first approach for claim importance ranking, adaptive credibility filtering for source quality, and a three-level parallelization architecture for latency optimization.

Evaluation on 30 annotated YouTube videos across health, climate, and political topics demonstrates strong performance: 81.3% claim extraction precision, 94.7% evidence retrieval success rate, and 73.3% verdict accuracy. The system achieves practical efficiency with mean processing time of 129.6 seconds per video at \$0.003 average cost, enabling cost-effective deployment for individual users.

The complete system, including a web-based interface with real-time streaming, is publicly available as open-source software, bridging the gap between academic research and practical fact-checking tools accessible to non-expert users.

Keywords: automated fact-checking, large language models, multi-agent systems, natural language processing, misinformation detection, YouTube, information retrieval, evidence-based verification

Contents

Abstract	vii
Table of Contents	ix
List of Figures	xi
List of Tables	1
1 Introduction	3
1.1 Context and Motivation	3
1.2 Project Objectives	5
1.3 Ethical Considerations and Social Impact	7
1.4 Methodology	8
1.5 Summary of Project Outputs	9
1.6 Document Structure	9
2 State of the Art	10
2.1 Multi-Agent Architectures for Fact-Checking	10
2.2 Automatic Claim Detection in Text	11
2.3 Retrieval-Augmented Generation	12
2.4 LLM-Based Claim Verification Methods	14
2.5 Evaluation of Fact-Checking Systems	16
2.6 Current Limitations and Research Opportunities	18
2.7 Positioning of This Work	19
3 Materials and Methods	21
3.1 System Architecture Overview	21
3.2 Technology Stack	24
3.3 Pipeline Components	25
3.4 Latency Optimization Strategies	33
3.5 Experimentation and Evaluation Framework	35
3.6 API Layer and User Interface	37

3.7	Engineering Practices	38
4	Results	39
4.1	Experimental Setup	39
4.2	Precision-Recall Tradeoff Analysis	41
4.3	Claim Extraction Performance	43
4.4	Verdict Generation Performance	44
4.5	Evidence Retrieval Performance	46
4.6	System Efficiency	47
4.7	Qualitative Analysis	48
4.8	Considerations for Generative AI Systems	50
4.9	Summary of Key Findings	52
5	Conclusions and Future Work	52
5.1	Summary of Contributions	53
5.2	Achievement of Objectives	54
5.3	Critical Assessment of Methodology	55
5.4	Reflection on Ethical and Sustainability Considerations	56
5.5	Lessons Learned	57
5.6	Future Work	58
5.7	Concluding Remarks	60
6	Glossary	61
6.1	Terms	61
6.2	Acronyms	62
	Bibliography	63
	Appendices	68
A	Ground Truth Annotation Dataset	68
A.1	Video Corpus Summary	68
A.2	Annotation Schema and Guidelines	70
A.3	Dataset Statistics by Category	71
A.4	Importance Score Distribution	71
A.5	Expected Verdict Distribution	71
A.6	Data Availability	71

List of Figures

1	High-level pipeline architecture showing the five main stages and three parallelization levels. Parallelization occurs at claims (Level 1), queries per claim (Level 2), and search results per query (Level 3). Verdict generation happens within Level 1 after each claim’s evidence is collected.	23
2	Online Search pipeline showing the four sequential steps executed for each search result. Steps 2–4 run in parallel across all K results per query (Level 3 parallelization).	30
3	User interface screenshots showing the fact-checking workflow: (a) users paste a YouTube URL, (b) real-time progress is displayed during analysis, (c) results are presented alongside the embedded video, and (d) each claim shows its verdict, confidence level, and supporting evidence with source reliability ratings.	38
4	Precision-recall curve for different <code>max_claims</code> values. The selected operating point ($k = 5$) achieves 81.3% precision at 26.2% recall, balancing claim quality with coverage.	42

List of Tables

1	Core technology stack	24
2	LLM providers and pricing	25
3	Claim importance scoring guidelines	27
4	Query type taxonomy	28
5	Evidence quality score breakdown	33
6	Operations using classical methods	35
7	Evaluation dataset statistics	40
8	Precision-recall tradeoff across different <code>max_claims</code> configurations	42
9	Claim extraction performance (n=30 videos, <code>max_claims</code> =5)	43
10	Verdict generation performance (n=30 videos)	44
11	Verdict accuracy distribution with evidence retrieval rates	45
12	Evidence retrieval performance (n=30 videos)	46
13	Source reliability distribution (n=724 total sources)	46
14	System latency (n=30 videos)	47
15	Examples of successful claim extraction with semantic matching	48
16	Verdict error categories	49
17	Low-accuracy video analysis (<25% verdict accuracy)	50
18	Evaluation video corpus by category and content type	69
19	Ground truth statistics by category	71
20	Distribution of importance scores across ground truth claims	71
21	Distribution of expected verdicts across ground truth claims	71

1 Introduction

Misinformation on digital platforms has become a major challenge in the information age. YouTube, with over 2 billion monthly active users, has evolved into a primary information source for millions of people worldwide. However, the absence of effective verification mechanisms allows false or misleading claims to propagate at massive scale, impacting public health, democratic processes, and social cohesion [2, 35].

This thesis addresses the critical need for automated fact-checking systems capable of processing video content from platforms like YouTube. The project develops Factible, a multi-agent system that implements an end-to-end pipeline for extracting, verifying, and synthesizing factual claims from YouTube videos using large language models (LLMs) and web-based evidence retrieval.

1.1 Context and Motivation

1.1.1 The Problem of Video Misinformation

The central problem this project addresses is the **lack of automated, accessible mechanisms to verify factual claims in audiovisual content from digital platforms**, specifically YouTube videos. Individual users lack tools to systematically verify claims, while manual verification requires advanced skills, considerable time, and resources that are not always available.

YouTube presents unique challenges for fact-checking:

- **Scale:** Over 500 hours of video are uploaded every minute, making human-only verification impossible.
- **Consumption patterns:** Videos are consumed passively, with viewers less likely to critically evaluate claims than when reading text.
- **Lack of annotation:** Unlike text articles, videos do not come with metadata identifying factual claims.
- **Temporal context:** Claims are embedded in a narrative flow, making extraction and verification more complex than isolated text statements.

1.1.2 Relevance and Importance

This project is relevant for multiple reasons:

- **Social impact:** Misinformation erodes trust in institutions, influences political decisions, and affects public health [2, 35]. An automated verification system can contribute to mitigating these effects by enabling faster, more accessible fact-checking.
- **Technological challenge:** The project integrates multiple areas of data science research: natural language processing, information retrieval, source credibility evaluation, and automated reasoning. It represents a technically complex problem requiring innovative solutions.
- **Scalability necessity:** The volume of content generated daily makes purely human fact-checking infeasible. Automation is necessary to address the problem at scale, serving as a force multiplier for human fact-checkers.
- **Media literacy:** Tools that empower citizens to critically evaluate information promote critical thinking and strengthen media literacy in society, contributing to a healthier information ecosystem.

1.1.3 Current State of Solutions

Currently, three main approaches exist for fact-checking:

1. **Manual fact-checking:** Organizations like Newtral, Maldita.es, PolitiFact, and FactCheck.org employ specialized journalists. This approach guarantees quality but is limited by human resources and is inherently slow, typically checking only a few claims per day.
2. **Academic systems:** Projects like FEVER [32], FEVEROUS [1], and recent work on scientific verification [34] operate primarily on structured knowledge bases (Wikipedia) or require specifically curated datasets. They are not available as products usable by end users and have limited scope.
3. **Verification APIs:** Google Fact Check Tools API aggregates fact-checks already published by professional organizations but only covers content previously verified manually. It does not provide new verification for unexamined content.

Identified gap: None of these approaches provides automated end-to-end verification for continuously generated audiovisual content on platforms like YouTube. A gap exists between academic research (static datasets, limited domains) and practical, accessible tools for end users.

1.2 Project Objectives

1.2.1 Main Objective

Design, implement, and evaluate a multi-agent system based on language models capable of automatically verifying factual claims in YouTube videos through retrieval and analysis of web-based evidence, generating substantiated verdicts with confidence estimates that are useful for end users.

1.2.2 Specific Objectives

1. Multi-agent architecture design

- Define a modular processing pipeline with specialized components
- Establish structured data schemas for communication between agents
- Design the information orchestration system

2. System component implementation

- Develop the module for extracting video transcripts from YouTube
- Implement the claim identification agent with relevance classification
- Create the optimized search query generator
- Develop the web scraping system with source reliability evaluation
- Implement the generator for explainable and traceable verdicts

3. LLM usage optimization

- Compare performance between commercial models (e.g., GPT-4o-mini) and open-source alternatives (e.g., Qwen 3, DeepSeek)
- Analyze trade-offs between cost, latency, and output quality
- Evaluate prompt optimization strategies and structured outputs

4. System evaluation

- Define technical performance metrics: response time, cost per token, throughput, and completion rate per agent
- Include LLM-specific quantitative metrics such as faithfulness, factual consistency, coherence, and stance accuracy

- Incorporate quantitative and qualitative evaluation with LLM-as-a-judge and manual review
- Conduct case studies across different video topics and analyze system limitations

5. End-to-end system implementation

- Develop a REST API with real-time streaming
- Create an intuitive web interface for end users
- Ensure reproducibility and generate documentation

6. Ethical and bias considerations

- Analyze potential biases in sources and models used
- Identify limitations of the automated approach, its possible impact on misinformation propagation, and define best practices to mitigate these risks during verification

1.2.3 Project Scope

Includes:

- YouTube videos with available transcripts (automatic or manual subtitles)
- Factual claims verifiable through public web sources
- Content in English
- On-demand processing of individual videos

Does not include:

- Self-generation of transcripts from audio (speech-to-text)
- Verification of visual content
- Large-scale batch verification
- Content from platforms other than YouTube

1.3 Ethical Considerations and Social Impact

1.3.1 Ethical Commitment and Global Competence

The project addresses three dimensions of ethical and global competence:

(I) Sustainability: The use of LLMs implies significant energy consumption. This is addressed through evaluation of local open-source models versus commercial cloud services, prompt optimization to minimize processed tokens, and transparent documentation of resource consumption.

(II) Ethical behavior and social responsibility: LLMs can perpetuate biases present in their training data. Mitigation strategies include:

- Total transparency in consulted sources with verifiable URLs
- Presentation of evidence showing different stances (supports/refutes/neutral)
- Explicit evaluation of confidence level in each verdict
- Clear warnings about limitations of the automated system

The system positions itself as a support tool that automates preliminary steps, allowing professionals to focus on complex analyses requiring human judgment.

(III) Diversity and human rights: The system seeks to evaluate diversity of perspectives when multiple web sources with different approaches are available, documenting this limitation when not possible. It does not censor content but provides additional context through verification, respecting freedom of expression. It does not process personal data or identifiable user information.

1.3.2 Sustainable Development Goals (SDGs)

The project contributes to three SDGs from the 2030 Agenda:

SDG 4 (Quality Education): The system acts as a media literacy tool that promotes critical thinking and the ability to evaluate information sources. It facilitates learning about fact-checking, evidence analysis, and recognition of verifiable claims—fundamental competencies in today's digital society.

SDG 16 (Peace, Justice and Strong Institutions): It combats misinformation that erodes trust in democratic institutions and media. It provides transparent, traceable, and accessible verification mechanisms that allow citizens to systematically contrast factual claims, contributing to a healthier, more resilient information ecosystem.

SDG 10 (Reduced Inequalities): It democratizes access to fact-checking through a free, open-source tool, reducing the gap between professional organizations with specialized resources

(fact-checkers, journalists) and individual citizens. It empowers users without specialized training to critically evaluate content on digital platforms.

1.4 Methodology

1.4.1 Research Strategy

This project adopts a **Design and Creation** strategy [25], appropriate for information systems research where a new technological artifact is developed whose construction and evaluation constitute the main contribution to knowledge. According to Hevner et al. [12], this strategy requires: (1) creating a viable artifact, (2) addressing a relevant problem, (3) conducting rigorous evaluation, (4) making clear contributions, (5) applying rigor in construction, (6) designing as a search process, and (7) effectively communicating results.

Data collection techniques for evaluation:

Quantitative:

- Automated measurement of system metrics: processing time, cost ($\text{tokens} \times \text{price}$), number of retrieved sources, stance distribution
- Evaluation using LLM-as-a-judge to assess quality of extracted claims and generated verdicts
- Multi-agent coordination metrics: task completion rate, latency per agent

Qualitative:

- Manual review of system outputs (claim relevance, verdict coherence)
- Detailed case studies across different topics

1.4.2 Development Methodology

Chosen strategy: Development of a new end-to-end product with a functional web interface.

Justification: This strategy is most appropriate because it allows evaluation of the complete practical viability of the system under real usage conditions, not just as a theoretical concept. Modular development from scratch facilitates experimentation with different LLM configurations and guarantees total control over the verification pipeline. Including a web interface demonstrates applicability for end users and allows collecting feedback on real system usability.

Iterative development process:

1. Implementation of individual components with isolated testing

2. End-to-end pipeline integration (functional MVP)
3. Iterative refinement module by module based on results
4. Comparative evaluation with different LLM configurations

1.5 Summary of Project Outputs

This thesis produces the following outputs:

1. **Factible**: A complete, open-source multi-agent fact-checking system for YouTube videos, available at <https://github.com/begoechavarren/factible>
2. **Modular pipeline architecture**: Five specialized components (Transcriptor, Claim Extractor, Query Generator, Online Search, Output Generator) with well-defined interfaces
3. **Web interface**: React-based frontend with real-time SSE streaming for interactive verification
4. **REST API**: FastAPI-based backend enabling programmatic access to fact-checking functionality
5. **Evaluation framework**: Infrastructure for systematic experimentation including tracking, metrics computation, and result analysis
6. **Annotated evaluation dataset**: 30 YouTube videos with 503 ground truth claims across health, climate, and political topics
7. **This thesis document**: Comprehensive documentation of the design, implementation, and evaluation of the system

1.6 Document Structure

The remainder of this thesis is organized as follows:

- **Section 2 (State of the Art)**: Reviews existing literature on automated fact-checking, multi-agent architectures, claim detection, retrieval-augmented generation, and LLM-based verification methods. Identifies gaps that this work addresses.
- **Section 3 (Materials and Methods)**: Presents the complete technical implementation of Factible, including system architecture, technology stack, pipeline components, latency optimization strategies, and the experimentation framework.

- **Section 4 (Results):** Reports experimental evaluation results across 30 annotated videos, including claim extraction performance, verdict accuracy, evidence retrieval statistics, and qualitative analysis.
- **Section 5 (Conclusions and Future Work):** Summarizes key findings, discusses achievement of objectives, reflects on ethical considerations, and proposes directions for future research.
- **Glossary:** Defines key terms and acronyms used throughout the document.
- **Appendices:** Provides supplementary material including the complete ground truth annotation dataset.

2 State of the Art

Automated fact-checking has evolved significantly since the mid-2010s, from rule-based systems to architectures using Large Language Models (LLMs). This evolution has been driven by advances in natural language processing, retrieval-augmented generation (RAG), and multi-agent systems that break down fact-checking into specialized, coordinated tasks.

This section examines automated fact-checking systems, emphasizing multi-agent architectures for video content verification. It analyzes the technological foundations of each pipeline component, from claim detection to evidence retrieval to verdict synthesis, and identifies gaps limiting practical deployment. The review covers literature from 2017 to 2025, focusing on systems combining LLMs with information retrieval and structured reasoning.

2.1 Multi-Agent Architectures for Fact-Checking

2.1.1 Foundational Multi-Agent Frameworks

FactAgent [38] introduced an agentic workflow that explicitly emulates the methodology of human fact-checkers. Instead of fine-tuning models for specific tasks, FactAgent employs a pre-trained LLM that operates through a structured script: (1) gathering evidence via tools, (2) analysing that evidence, and (3) synthesising a verdict. A key advantage of this approach is its zero-shot nature: the system uses the LLM’s existing capabilities without requiring task-specific training data. However, the sequential multi-step process can introduce latency, and errors in early stages may cascade through the pipeline.

LoCal [3] handles complex claims through a decomposition–reasoning–evaluation loop. It uses specialized agents: a decomposer breaking claims into sub-claims, reasoning agents verifying each with evidence, and evaluators ensuring logical consistency and testing counterfactual

scenarios. If inconsistencies arise, agents iteratively revise their reasoning. This targets single-pass verification weaknesses but increases computational cost.

Multi-agent debate systems [14, 19] use adversarial collaboration where distinct LLM agents take opposing roles: one argues for a claim’s truth, another against it, while a judge decides the outcome. This exposes contradictions and forces agents to defend positions with evidence, reducing confirmation bias and hallucinations. The adversarial setup prevents premature convergence on incorrect conclusions, though debates can reach impasses and quality depends on the judge agent’s synthesis ability.

2.1.2 Evolution and Current Landscape

Multi-agent design for fact-checking emerged recently (2023–2025). Early research used linear pipelines where a single model performed one task at a time. By 2024, works like FactAgent and LoCal began breaking verification into specialized agents [3, 38]. In 2025, researchers added debate mechanisms, self-reflection, and richer tool use [13, 19, 33]. MAD-Sherlock showed that debate-driven systems reduce hallucinations through collaborative verification [14]. Despite progress, recent evaluations reveal latency challenges, high compute requirements, and the need for careful orchestration to avoid loops or premature termination.

2.2 Automatic Claim Detection in Text

Identifying which statements need fact-checking is crucial for any verification pipeline. For YouTube videos, this means scanning transcripts to flag factual claims that are verifiable and important enough to check. This task is known as claim check-worthiness detection.

2.2.1 Traditional Supervised Approaches

Claim detection research began in the mid-2010s, targeting political speech. **ClaimBuster** was pioneering work, the “first-ever end-to-end fact-checking system” [10]. It used a supervised model trained on human-labelled debate transcripts to score sentences for verifiable factual claims. Using feature engineering and early neural networks, it achieved sufficient accuracy for integration by fact-checking organizations like Duke Reporters’ Lab [10]. ClaimBuster automated the initial triage step, helping journalists prioritize statements from debates or speeches.

Subsequent work refined datasets and models. **ClaimRank** expanded U.S. presidential debate data and introduced context-aware modeling, using surrounding sentences to improve detection [7]. The **CLEF CheckThat!** lab (2018–2022) released annual challenges with social media posts or political statements in multiple languages, labelled for check-worthiness [32].

These efforts established claim detection as a classification or ranking problem, with BERT and transformers becoming dominant after 2018. The key insight was that check-worthiness requires assessing both **importance** and **verifiability**. “The sky is blue” is factual but not check-worthy due to obviousness. “Unemployment rose by 15% last quarter” is both factual and check-worthy because it’s verifiable and important. Training models to capture this required carefully annotated datasets with clear guidelines.

2.2.2 LLM-Based Claim Detection

Recent work explores whether large language models can identify check-worthy claims without fine-tuning. **Sawinski et al. (2024)** compared fine-tuned BERT variants with GPT-3/GPT-4 in zero-shot or few-shot mode [29]. Simple zero-shot LLM prompts still underperform fine-tuned models on benchmarks. LLMs often have inconsistent internal definitions of “worthiness” and are sensitive to prompt wording, while fine-tuned models have learned explicit criteria from labelled data.

However, careful prompts can substantially improve LLM performance. **Li et al. (2023)** built a fully automated fact-checking prototype with an LLM-based claim detection module, using GPT-3 with verbose few-shot prompts [17]. While quantitative metrics weren’t reported, the work demonstrated that LLMs can effectively replace traditional claim detectors in automated pipelines.

Ni et al. (2024) proposed a three-step prompting approach for consistent claim identification [23]. The LLM analyzes text in stages: highlighting factual statements, applying check-worthiness criteria, and ranking by importance, similar to chain-of-thought reasoning. This improved consistency but focused on verifiable claim identification rather than worthiness ranking [24].

2.2.3 Current Challenges

Key challenges include: (1) **definition ambiguity**: what constitutes a “check-worthy” claim varies by context; (2) **scalability**: scanning long transcripts with LLMs is slow and expensive; (3) **false positives**: overly aggressive detection wastes resources; and (4) **domain adaptation**: models trained on political debates may not work for scientific or economic content.

2.3 Retrieval-Augmented Generation

A core pillar for automated fact-checking is retrieval-augmented generation (RAG), which combines text generation with external information retrieval to ensure outputs are based on verifiable sources.

2.3.1 RAG Fundamentals

Lewis et al. (2020) formalized RAG by showing that augmenting generation with external knowledge retrieval significantly improves performance on knowledge-intensive tasks [15]. For fact-checking, RAG is essential: systems must fetch reliable sources that support or refute claims. The **FEVER** dataset exemplified this retrieve-then-verify pattern: given a claim, retrieve relevant documents (e.g., Wikipedia pages), then determine if they support or refute the claim [32]. Modern systems use LLMs for both retrieval and verification, conditioning their reasoning on retrieved evidence.

RAG addresses a critical limitation of pure LLM approaches: parametric knowledge can be outdated, incomplete, or hallucinated. By retrieving external information (from document indexes, web search, or APIs), RAG systems access current information, provide source attribution, and ground reasoning in verifiable evidence. This is crucial for fact-checking, where claims often reference recent events, specific statistics, or specialized knowledge not in LLM training data.

2.3.2 Tool Use and Web Retrieval

Integrating tool use with LLMs has enabled more sophisticated retrieval. The **ReAct pattern** (Reason and Act) interleaves tool use with chain-of-thought reasoning, letting LLMs decide when to call external tools like search engines [37]. **WebGPT** demonstrated training LLMs to use web browsers to answer questions and cite sources, improving factual accuracy by teaching models to navigate search results and synthesize information from multiple pages [22]. Similarly, **Toolformer** showed that LLMs can be fine-tuned to call external tools like search engines or calculators, reducing hallucinations by grounding answers in retrieved evidence [30].

Recent work has applied these techniques to fact-checking. **Chern et al. (2023)** proposed using Google Search, Google Scholar, and other tools to verify LLM-generated text against external sources [4], while **Cheung and Lam (2023)** combined search-engine retrieval with LLaMA to predict claim veracity [5]. These tool-augmented methods address LLMs' inherent knowledge limitations, which can be outdated or incomplete [4].

2.3.3 Open Web versus Closed Knowledge Bases

Most academic fact-checking systems restrict retrieval to trusted corpora (primarily Wikipedia) to simplify evaluation and ensure evidence quality. This yields high precision but severely limits real-world coverage [6]. The FEVER dataset, while influential, exemplifies this by assuming all claims can be checked against a June 2017 Wikipedia snapshot, which breaks down for recent events, specialized domains, or claims requiring sources like World Bank reports or CDC

guidelines.

Tian et al. (2024) integrated web-retrieval agents into an LLM pipeline and demonstrated improved misinformation detection [33]. However, open-web retrieval introduces challenges: (1) **source credibility**: not all websites are reliable; (2) **information quality**: web content varies in accuracy; (3) **ranking complexity**: identifying relevant sources among millions of candidates; and (4) **dynamic nature**: content changes, affecting reproducibility.

Current best practices include prioritizing sources with high domain authority (established news organizations, academic institutions, government agencies), cross-referencing multiple independent sources, explicitly evaluating source credibility using metadata (publication date, author credentials, institutional affiliation), and maintaining transparency by exposing retrieved sources to users. Systems like FactAgent incorporate evidence retrieval as a dedicated step, using search tools to query the web and filtering results based on relevance and credibility [38].

2.3.4 Query Optimization for Fact-Checking

A critical but often overlooked component of RAG systems is query formulation. The same claim can be verified or refuted depending on how search queries are constructed. Query quality significantly impacts downstream task performance [10, 19].

Effective query optimization involves several strategies. **Keyword extraction** identifies salient terms likely to appear in relevant sources, filtering stop words and focusing on entities and key concepts. **Query expansion** generates multiple variants to capture different phrasings; for example, expanding “unemployment rate increased” to also search for “jobless claims rose” or “labour market deterioration” [15]. **Entity recognition** identifies named entities (people, organizations, locations) that should be included in queries, as these serve as strong signals for retrieval. **Temporal awareness** incorporates time constraints when claims reference specific periods, adding the relevant year when verifying recent events.

Recent multi-agent systems often dedicate a specialized agent to query generation, recognizing this step significantly impacts retrieved evidence quality [38]. Poor queries may miss relevant sources or retrieve irrelevant information, degrading overall performance regardless of verification model quality. FactAgent includes explicit query formulation as one of its agent steps, using the LLM to generate search-optimized queries [38].

2.4 LLM-Based Claim Verification Methods

Once claims are identified and evidence retrieved, systems must determine veracity, labelling claims as supported (true), refuted (false), or not enough evidence. Traditional approaches treated this as textual entailment, using neural classifiers to determine if evidence entails or

contradicts claims. With LLMs, a new approach emerged: using models to perform verification through natural language reasoning.

2.4.1 Prompting Strategies

Zero-shot and few-shot prompting involves providing an LLM with a claim and evidence, asking it to decide veracity and explain why: “Claim: X. Evidence: [text]. Based on the evidence, is the claim true or false?” [39]. In zero-shot mode, the LLM relies on internal reasoning and evidence interpretation. In few-shot mode, the prompt includes examples of claims with evidence and the correct verdict to guide the model.

GPT-4 and similar models show surprising capability at this task, often correctly interpreting whether evidence supports statements. However, LLMs can be overly agreeable, sometimes hallucinating justifications or defaulting to “Supported” even when evidence is insufficient [39]. This confirmation bias stems from models’ training to be helpful and provide answers, even when saying “I don’t know” would be more appropriate.

Careful prompt engineering can help. Effective strategies include explicitly instructing the model to answer “Not Enough Evidence” when information is insufficient, adding system messages emphasizing accuracy over helpfulness, requesting citation of specific evidence sentences supporting its verdict, using temperature settings near zero to reduce randomness, and implementing multi-pass verification where the model first generates a verdict then critiques its own reasoning.

2.4.2 Chain-of-Thought Reasoning

More advanced approaches use **chain-of-thought reasoning** or implement the LLM as an agent in a loop. The **ReAct pattern** has the LLM explicitly reason step-by-step while using tools [37]. For verification, this might involve: (1) breaking the claim into parts, (2) querying a search engine for each part, (3) evaluating each piece of evidence, and (4) synthesizing a conclusion.

FactAgent’s structured workflow exemplifies this: the LLM follows a script where each step (search, read results, extract evidence, cross-check, formulate verdict) is explicit and logged for transparency [38]. This approach is flexible; if initial evidence is inconclusive, the agent can trigger refined searches. Also, the zero-shot operation without training, mimics human fact-checker processes [38].

Chain-of-thought provides transparency (each reasoning step is explicit and inspectable), debuggability (identifying which step failed), and explainability (the reasoning trace serves as a natural language explanation). However, multiple LLM calls for each sub-step can be slow and

expensive, and errors compound across stages. If claim decomposition is incorrect, subsequent reasoning will be compromised regardless of retrieval and evaluation quality.

2.4.3 Self-Consistency and Verification

SelfCheckGPT introduced self-consistency checking for hallucination detection by generating multiple independent answers and checking if factual claims agree across responses [20]. The approach operates on the principle that hallucinated information varies across samples while grounded information remains consistent, allowing systems to identify low-confidence or potentially hallucinated components.

LLM-as-a-judge approaches use one model to generate answers and another to verify them, catching errors before presenting results [27, 28]. **Cross-model checking** extends this by using different models to verify outputs; if they disagree, the system can abstain or reconsider [1]. Many verification pipelines also use **stance detection** models to classify evidence snippets as supporting, refuting, or not mentioning the claim [32]. While these verification techniques improve reliability, they add computational overhead through additional model calls and processing steps.

2.4.4 Current Capabilities and Limitations

Carefully prompted LLMs can achieve near state-of-the-art performance on tasks like FEVER [32]. However, they still make mistakes, especially on ambiguous or complex claims requiring specialized knowledge or multi-step reasoning.

A noted limitation is the tendency to default to “Supported”: models sometimes erroneously agree claims are true if any related evidence is found (confirmation bias), rather than truly verifying the exact claim [39]. Designing prompts or agent behaviors to be appropriately skeptical and output “Not Enough Info” when evidence is lacking remains important, requiring calibration to avoid being too trusting or too skeptical.

2.5 Evaluation of Fact-Checking Systems

Evaluating automated fact-checking systems requires assessing accuracy, explanation quality, evidence usage, and practical usability.

2.5.1 Veracity Classification Metrics

Standard classification metrics include accuracy, F1-score, and precision/recall. The FEVER challenge introduced the **FEVER score**, which requires both correct labels and proper evidence, penalizing systems that get labels right without grounding [32]. For multi-class truth

scales (e.g. PolitiFact’s “true” to “pants on fire”), accuracy within each class or Cohen’s kappa can be used.

2.5.2 Evidence Retrieval Metrics

Key metrics include **Recall@k** (whether correct evidence appears in the top k retrieved documents), **Precision** (proportion of relevant evidence), and **Mean Average Precision (MAP)** (accounting for both relevance and ranking order) [32]. End-to-end evaluations typically credit systems only when they retrieve human-identified evidence, though multiple valid sources may exist for a claim.

2.5.3 Explanation Faithfulness and Quality

Explanations should be **faithful** (reflect actual reasoning) and **factually consistent** with evidence. While automatic metrics like BLEU or ROUGE exist, they don’t measure factuality well [27]. More sophisticated approaches include FactCC [31], Q², and entailment checks. **LLM-as-a-judge** has become popular, using LLMs to score explanation coherence and factuality, though these judges may have biases and require validation against human assessments [27, 28].

2.5.4 Logical Coherence and Consistency

Stance consistency checks whether evidence stances align with final verdicts: if a system claims truth but all evidence refutes it, that indicates errors. LoCal evaluates logical consistency by checking if composed solutions imply claim veracity [3]. Secondary models can perform entailment checks, creating a verification layer for the system’s reasoning.

2.5.5 Human Evaluation

Human judgment remains the gold standard. Researchers conduct user studies or expert assessments where fact-checking experts rate verdict correctness and reasoning soundness, while lay users evaluate whether explanations are convincing and understandable. Human evaluation also assesses readability, perceived trust, and completeness [19, 35].

2.5.6 Computational Performance Metrics

For practical deployment, computational efficiency matters. Relevant metrics include latency, throughput, monetary cost, and resource utilization. Although rarely reported in academic papers, these metrics are critical for operational systems targeting near-real-time analysis [33, 38].

2.6 Current Limitations and Research Opportunities

Despite rapid progress, automated fact-checking systems have significant limitations constraining practical deployment.

2.6.1 Lack of End-to-End Usability

Most research prototypes focus on isolated components rather than seamless end-to-end tools. Some excel at claim detection but assume manual verification [10], while others verify claims but require human identification. Even ClaimBuster, dubbed “end-to-end”, only highlighted claims without verification [10]. Complete systems need to integrate detection, verification, and source tracing, but existing systems typically address only one or two steps [18]. For YouTube videos, true end-to-end systems should handle transcription extraction, claim detection, evidence retrieval, verification, and user-friendly presentation, but this integration is rarely achieved [18].

2.6.2 Dependence on Structured Sources

Much research restricts evidence to structured knowledge bases, primarily Wikipedia. While this yields cleaner evaluation, it severely limits applicability [1]. Real misinformation often requires specialized sources not available in Wikipedia, and systems benchmarked on FEVER tend to be overfitted [32]. In practice, fact-checkers must handle the open web including news sites, scientific papers, and government databases, introducing challenges of source credibility and information quality [1].

2.6.3 Production Readiness Gap

Most solutions remain research-grade implementations rather than production-ready systems. Code is typically provided as research artifacts (Jupyter notebooks, command-line scripts) without the robustness, or user interfaces needed for public deployment [18]. Even public tools like Google Fact Check Explorer only search existing fact-checks rather than performing new verification [8]. Additionally, many advanced systems require computational resources not feasible without powerful hardware or costly API access, limiting their practical accessibility [18].

2.6.4 Performance and Scalability Limitations

Checking long videos with many claims poses significant challenges for any system. If verification takes minutes per claim, videos with 20 claims become impractical for interactive use.

Current research often doesn't address runtime performance [18]. Multi-agent systems can theoretically operate in parallel, but sequential dependencies limit parallelization. Cost is also a factor: using commercial LLMs like GPT-4 for every step can be prohibitively expensive at scale.

2.6.5 Trust and Transparency Issues

Users may be reluctant to trust AI verdicts without understanding their derivation, and many systems have been criticized as “black boxes” [35]. Multi-agent systems and chain-of-thought approaches attempt to address this via explicit reasoning traces, but face risks of hallucinated explanations or justifications. LLMs’ documented tendency to confidently present false information can undermine the system’s purpose in fact-checking contexts.

2.6.6 Lack of Comparative LLM Evaluation

Despite many LLM options (commercial models like GPT-4 or Claude; open-source models like LLaMA, Qwen, DeepSeek), there’s limited systematic comparison of their suitability for fact-checking tasks. Research tends to use whichever model is most accessible without careful comparison of trade-offs in cost, accuracy, and latency [27].

2.6.7 The Video Content Gap

Most fact-checking research focuses exclusively on text-based content, primarily analyzing written articles, social media posts, or pre-extracted claims from political debates [10, 32]. YouTube, despite having over 2 billion monthly active users and serving as a primary information source for millions, remains largely unaddressed by automated fact-checking systems. Existing research assumes pre-processed text, leaving the video-to-claim pipeline unsolved. Video content introduces unique challenges including automatic speech recognition errors, missing punctuation and formatting, temporal context loss, and high computational costs. Given YouTube’s central role in information consumption and its documented contribution to misinformation spread [2], this gap represents a critical limitation in automated fact-checking capabilities.

2.7 Positioning of This Work

This thesis addresses the critical gaps identified above by developing a complete multi-agent system for YouTube video fact-checking, making the following contributions.

2.7.1 End-to-End Video Fact-Checking System

The system implements a complete pipeline integrating all stages from transcription extraction to claim detection, query generation, evidence retrieval, and verdict synthesis, processing raw YouTube URLs autonomously. Unlike research prototypes assuming pre-processed inputs, this handles the full workflow from video to verified claims, addressing the end-to-end usability gap [18].

2.7.2 Open-Web Evidence Retrieval

Moving beyond Wikipedia and closed knowledge bases, the system retrieves evidence from the live web using search engines, incorporates source credibility evaluation, and checks results across multiple sources. This addresses the limitation of systems constrained to structured sources and enables verification of claims about recent events, specialized domains, and topics not well covered in encyclopedic sources [1, 32].

2.7.3 Practical User Interface

A web-based interface with real-time streaming of intermediate results using Server-Sent Events makes the verification process transparent and accessible to non-expert users. This bridges the gap between research prototypes and usable products [18], demonstrating that academic advances can be packaged for practical use.

2.7.4 Systematic LLM Comparison

The thesis conducts comparative evaluation of different LLM configurations (commercial versus open-source models, different model sizes), analyzing trade-offs in cost, latency, and quality across pipeline components. This fills a gap in the literature where such trade-offs are rarely studied explicitly, providing practical guidance for deployment decisions [27].

2.7.5 Modular Multi-Agent Architecture

The system implements five specialized agents with structured data schemas enabling transparency, maintainability, and future updates. Each agent can be evaluated and optimized independently, and the modular design facilitates experimentation with different models and techniques, showing the practical benefits of multi-agent approaches in a production-oriented context [3, 19, 38].

2.7.6 Comprehensive Evaluation Framework

The evaluation combines quantitative metrics (technical performance, LLM-specific metrics like faithfulness and consistency), LLM-as-a-judge evaluation, and qualitative case studies across different video topics. This multifaceted approach addresses the evaluation challenges discussed earlier and provides a realistic assessment of system capabilities and limitations [27, 28].

2.7.7 Focus on Video Platform

Unlike most fact-checking research that focuses on text-based content, this work specifically addresses YouTube video verification, handling the complete video-to-verification workflow including transcript extraction, temporal claim localization, and context-aware verification. While existing systems assume pre-processed text inputs [10, 32], this implementation processes raw video URLs end-to-end, filling a significant gap where YouTube has been largely overlooked by fact-checking research despite its scale and influence. This work demonstrates that automated fact-checking techniques can be successfully applied to video platforms, extending the scope of verification systems beyond traditional text-based media.

By combining state-of-the-art techniques from multi-agent systems, RAG, and LLM-based verification with explicit focus on practical deployment, this work advances automated fact-checking from research prototype toward usable tool.

3 Materials and Methods

This section presents the comprehensive technical implementation of Factible, a multi-agent system for automated fact-checking of YouTube videos. The complete source code is publicly available at <https://github.com/begoechavarren/factible>. The system implements an end-to-end pipeline that processes video content through five specialized components, leveraging large language models (LLMs) for reasoning tasks while employing classical algorithms for deterministic operations. The implementation follows design-science principles [12, 25] and focuses on building practical artifacts that are iteratively evaluated and refined.

3.1 System Architecture Overview

3.1.1 High-Level Architecture

Factible implements an end-to-end automated fact-checking pipeline for YouTube videos using a multi-agent architecture. Recent research on LLM agents demonstrates that multi-agent collaboration can enhance factuality and reasoning by allowing specialized agents to coordinate

on tasks [36]. FactAgent further shows that decomposing fact-checking into dedicated agents for input ingestion, query generation, evidence retrieval, and verdict prediction yields higher accuracy and transparency [38]. The Factible architecture follows this line of work by processing video content through five specialized, modular components that operate sequentially with three levels of internal parallelization.

The system processes a YouTube video URL through five sequential stages, each with specialized responsibilities. The pipeline begins with transcript extraction, proceeds through claim and query generation, conducts online evidence retrieval, and culminates in structured verdict synthesis. This modular design enables independent optimization of each component while maintaining clear data contracts between stages.

The five stages are:

1. **Transcriptor**: Extracts video transcripts via YouTube Transcript API, preserving timestamped segments for claim localization. The component includes automatic fallback to proxy service when rate-limited, ensuring robust transcript retrieval across different access conditions.
2. **Claim Extractor** (LLM Agent): Employs thesis-first reasoning to infer the video’s central argument before extracting factual and verifiable claims. Each claim receives an importance score based on its impact on the video’s thesis. Post-processing uses fuzzy string matching to locate claims within the transcript for timestamp mapping.
3. **Query Generator** (LLM Agent): Generates diverse search queries across four strategic types—direct, alternative, source-seeking, and contextual. Each query receives a priority score (1–5) based on evidence likelihood, enabling filtering of low-priority queries.
4. **Online Search**: Executes a four-step evidence retrieval pipeline for each query: (i) Google Search via Serper API, (ii) website reliability assessment using Media Bias/Fact Check (MBFC) data combined with domain heuristics [21], (iii) content fetching via Selenium WebDriver with JavaScript rendering support, and (iv) LLM-based evidence extraction with stance classification (supports, refutes, mixed, unclear).
5. **Output Generator** (LLM Agent): Synthesizes evidence into structured verdicts by building evidence bundles grouped by stance, generating natural language summaries with confidence levels, calculating algorithmic evidence quality scores, and mapping claims to video timestamps for interactive navigation.

Figure 1 illustrates the complete pipeline architecture with data flow and parallelization points across all five stages.




Figure 1: High-level pipeline architecture showing the five main stages and three parallelization levels. Parallelization occurs at claims (Level 1), queries per claim (Level 2), and search results per query (Level 3). Verdict generation happens within Level 1 after each claim’s evidence is collected.

3.1.2 Design Principles

The system adheres to several key design principles derived from software engineering best practices and GenAI application development [12]:

- Modularity**: Each component is isolated with well-defined inputs and outputs using Pydantic schemas, enabling independent optimization, testing, and replacement.
- Structured Outputs**: All LLM interactions use Pydantic AI with typed output schemas, ensuring type safety, automatic validation, and consistent data structures across the pipeline.
- Transparency**: The full evidence chain is preserved and exposed to users—sources, reliability ratings, stances, and reasoning are all traceable from final verdict back to original source.
- Progressive Enhancement**: The pipeline operates with graceful degradation (e.g., fall-back to snippet if scraping fails, fallback to proxy if rate-limited) rather than failing entirely.

5. **Cost-Conscious Design:** Configurable limits (`max_claims`, `max_queries`, `max_results`) prevent runaway API costs and control latency during development and production.
6. **Reproducibility:** Deterministic LLM outputs (`temperature=0.0`), structured YAML configurations, and comprehensive experiment tracking enable reproducible research.
7. **Separation of Concerns:** Classical algorithms handle tasks like reliability scoring, deduplication, and quality calculation, reserving LLM calls for tasks requiring reasoning and language understanding.

3.2 Technology Stack

3.2.1 Core Technologies

Table 1 presents the core technologies employed in the implementation.

Table 1: Core technology stack

Category	Technology	Version	Purpose
Language	Python	3.12	Core implementation with type hints
LLM Framework	Pydantic AI	$\geq 1.0.0$	Agent orchestration, structured outputs
Data Validation	Pydantic	$\geq 2.0.0$	Schema definitions, runtime validation
Web Framework	FastAPI	$\geq 0.115.0$	REST API with SSE streaming
HTTP Server	Uvicorn	$\geq 0.32.0$	High-performance ASGI server
Async HTTP	httpx	$\geq 0.28.1$	Async HTTP client
Web Scraping	Selenium	$\geq 4.15.2$	JavaScript-rendered content extraction
YouTube	youtube-transcript-api	$\geq 1.2.2$	Transcript extraction
Domain Info	python-whois	$\geq 0.8.0$	Domain age lookup
CLI	Typer	$\geq 0.15.0$	Experiment runner CLI
Analysis	pandas, matplotlib	-	Data analysis and visualization

3.2.2 Large Language Models

The system supports multiple LLM providers to enable comparison of cost-quality trade-offs. Table 2 shows the available models and their configurations.

Table 2: LLM providers and pricing

Provider	Model	Context	Pricing (per 1M tokens)	Use Case
OpenAI	gpt-4o-mini	128K	\$0.15 / \$0.60	Default
OpenAI	gpt-4o	128K	\$5.00 / \$15.00	High-quality
OpenAI	gpt-4-turbo	128K	\$10.00 / \$30.00	Premium
Ollama	qwen3:8b	40K	Free (local)	Budget/offline
Ollama	qwen3:4b	256K	Free (local)	Budget/offline

Large language models such as GPT-4 offer multimodal capabilities and demonstrate human-level performance across diverse benchmarks [26]. Despite these advances, models still suffer from hallucinations and are constrained by limited context windows, underscoring the need for careful configuration and reliability safeguards [26]. The model management layer therefore makes it straightforward to switch providers, tune limits, and keep outputs deterministic through shared configuration.

The same layer also exposes a zero-cost, offline alternative through Ollama with Qwen 3 variants (qwen3:8b, qwen3:4b). Switching between OpenAI and local providers requires only a configuration change, enabling experiments that balance latency, quality, and hardware constraints. GPT-4o-mini remains the default selection for this thesis because cloud inference provides lower latency than consumer hardware while maintaining strong reasoning quality.

3.2.3 External Services

The system integrates with external services for search and transcript extraction:

- **Serper API:** Google Search wrapper providing organic search results with approximately 2,500 queries per month on the free tier.
- **YouTube oEmbed API:** Video metadata retrieval without authentication.
- **Webshare Proxy:** Rate limit bypass for transcript extraction with configurable proxy locations.

3.3 Pipeline Components

3.3.1 Transcriptor

The Transcriptor component extracts YouTube video transcripts with precise timestamp information for later claim-to-video mapping.

Implementation Details The transcriptor uses the `youtube-transcript-api` library to fetch available transcripts, with preference for English. When rate-limited by YouTube, it automatically falls back to a proxy service (Webshare). Key features include:

- **Timestamped Segments:** Each segment preserves `start` time and `duration` in seconds.
- **Character Position Mapping:** Enables mapping claim text positions back to video timestamps.
- **Title Fetching:** Uses YouTube oEmbed API to retrieve video title for context.
- **Proxy Fallback:** Automatic retry through Webshare proxy when rate-limited.

Output and Timestamp Mapping Each run returns the video identifier, the full transcript text, and a list of timestamped segments containing text, start time, and duration. A lightweight mapping pass scans the segments sequentially while maintaining a cumulative character counter so that claim positions can be translated into timestamps. This enables the interface to jump straight to the moment where a claim appeared without duplicating transcript parsing logic elsewhere.

3.3.2 Claim Extractor

The Claim Extractor identifies factual, verifiable claims from video transcripts using LLM-based extraction with thesis-relative importance ranking. This approach builds on prior work in automated claim detection: supervised models trained on annotated political debates have been used to detect check-worthy claims [9], and end-to-end systems like ClaimBuster monitor public discourse and prioritize factual statements for manual fact-checking [16]. These systems show that focusing on salient, verifiable claims improves the efficiency of fact-checking pipelines.

LLM Configuration The claim extractor uses deterministic settings for reproducibility: temperature set to 0.0 to ensure consistent outputs across multiple runs, max tokens limited to 1,200 to control response length and latency, and automatic retry logic (3 attempts) to handle transient LLM failures gracefully.

Prompt Engineering Strategy: Thesis-First Approach The claim extractor employs a novel *thesis-first approach* with multi-step reasoning designed to prioritize claims most critical to the video’s central argument:

Step 1: Thesis Inference — Before listing claims, the LLM infers the video’s central thesis in no more than 25 words (e.g., “Climate change alarmism is driven more by politics and media than by settled science”).

Step 2: Importance Ranking with Thesis Impact Test — Claims are scored based on their impact on the video’s thesis using the question: “If this claim were proven false, would the thesis collapse or materially weaken?” Table 3 presents the scoring guidelines.

Table 3: Claim importance scoring guidelines

Score Range	Description	Examples
0.85–1.0	Prescriptive/causal claims undermining thesis	Policy proposals, causal mechanisms
0.60–0.80	Quantitative/historical evidence tied to thesis	Statistics, dates, expert citations
0.30–0.55	Context/supporting background	Definitions, general facts
0.0–0.25	Peripheral/anecdotal details	Personal stories, credentials

Step 3: Relevance Guardrails — Pure credential facts are capped at 0.30 unless the thesis questions expertise; statements not affecting the thesis are capped at 0.25; pure opinions are excluded; and paraphrases and duplicate numbers are removed.

Prompt Engineering Techniques The system employs several established prompting techniques across its components:

- **Role/system prompting:** Each component receives a focused system prompt that spells out its role, constraints, and expected outputs for consistent behavior.
- **Structured outputs:** Responses are constrained to typed schemas so downstream stages always receive validated fields.
- **Chain-of-thought reasoning:** The thesis-first steps require the model to state the thesis and then rate claims against it, which improves prioritization.
- **Dynamic instruction injection:** Runtime parameters such as `max_claims` are inserted directly into prompts so component behavior adapts without duplicating templates.
- **Deterministic generation:** Temperature stays at 0.0 across the pipeline for reproducible, debuggable outputs.
- **Zero-shot prompting:** Components receive task instructions without example I/O pairs, minimizing prompt length and avoiding format bias.

Post-Processing: Fuzzy Claim Localization After LLM extraction, each claim is located in the original transcript using fuzzy string matching. The algorithm normalizes both the claim text and transcript, then applies a sliding window approach (with ± 2 words tolerance around the expected claim length) to find the best matching region. Similarity is computed using Python’s sequence matching algorithm, which calculates the ratio of matching characters between two strings. A minimum similarity score of 0.5 is required for a valid match. This process yields the character-level start and end positions within the transcript, along with the match confidence score, enabling precise timestamp mapping from transcript positions back to video timestamps.

3.3.3 Query Generator

The Query Generator produces diverse, prioritized search queries tailored to each claim. A single LLM call emits direct restatements, synonym-heavy variations, source-focused prompts, and broader context queries together with priority tags so downstream steps can respect latency budgets. Guardrails ensure the suggestions include relevant entities, time periods, and qualifiers while capping each claim at a manageable number of high-yield searches.

Query Type Taxonomy The system generates four types of queries with different search strategies, as shown in Table 4.

Table 4: Query type taxonomy

Type	Description	Strategy	Example
DIRECT	Exact claim phrasing	Verbatim search	“unemployment rose 15% Q3 2024”
ALTERNATIVE	Rephrased with synonyms	Semantic variation	“jobless rate increase third quarter”
SOURCE	Target authoritative sources	Source-seeking	“BLS unemployment statistics Q3”
CONTEXT	Broader context	Background search	“economic indicators fall 2024”

Priority System Queries are prioritized 1–5 based on likelihood of finding reliable, definitive information: priority 1 queries are always included, priority 2 by default, priority 3 if budget allows, and priorities 4–5 rarely or only for completeness.

Context-Aware Query Generation Beyond factual accuracy, the query generator is designed to detect misleadingly framed claims—statements that may be technically accurate but presented without essential context. The system prompt instructs the LLM to respect temporal context when choosing keywords (e.g., including relevant years or qualifiers to avoid mixing eras), and to explicitly generate queries seeking counter-arguments or opposing views. This approach helps surface evidence that may qualify, limit, or contextualize the original claim, enabling more nuanced verdict generation.

3.3.4 Online Search

The Online Search component implements a multi-step pipeline to retrieve, assess, and extract evidence from web sources with adaptive quality filtering. Unlike the other LLM-based components, Online Search orchestrates multiple classical algorithms alongside a single LLM call for evidence extraction. This hybrid approach balances speed, reliability, and reasoning capabilities.

The reliability assessment combines domain-level heuristics with the Media Bias/Fact Check (MBFC) methodology, which employs a comprehensive weighted scoring system to evaluate media outlets’ ideological bias and factual reliability [21]. These scores, combined with stance-aware evidence extraction instructions, keep unreliable or speculative passages from flowing downstream.

Figure 2 illustrates the four-step Online Search pipeline executed for each query.




Figure 2: Online Search pipeline showing the four sequential steps executed for each search result. Steps 2–4 run in parallel across all K results per query (Level 3 parallelization).

Step 1: Google Search (Serper API) The Google search client wraps the Serper API for asynchronous search execution. The implementation uses persistent HTTP connections for performance and returns structured search hits containing title, URL, and snippet fields. Each query can retrieve up to 10 results, with the limit parameter controlling the exact number returned. The async design enables parallel query execution across multiple claims simultaneously.

Step 2: Website Reliability Assessment The reliability checker uses a multi-factor scoring system with first-match priority, combining external datasets with algorithmic heuristics:

- **Media Bias/Fact Check (MBFC) Dataset**: The system loads timestamped JSON snapshots containing nearly 10,000 news sources with credibility ratings (dataset extracted December 2025). Credibility mappings are: high → 0.85, medium → 0.60, low → 0.30, very low → 0.15.
- **TLD Reputation**: High-trust top-level domains (.gov, .edu, .int) receive a base score of 0.90, reflecting their institutional authority.
- **Domain Age via WHOIS**: Domains ≥ 10 years old receive a +0.10 bonus (established

presence), while domains <1 year old receive a -0.15 penalty (recent creation may indicate lower trust).

The output includes a categorical rating (high, medium, low, unknown), numerical score (0.0–1.0), reasoning for the assessment, and political bias classification when available from MBFC data.

Step 3: Content Fetching (Selenium) Content fetching uses headless Chrome with smart waits: the scraper first grabs paragraph elements immediately, then allows extra rendering time only when less than 100 characters were captured. Images are disabled, page-load and wait timeouts cap the work (20 seconds and 12 seconds, respectively), and the cleaned text is trimmed to 8,000 characters before being passed to the LLM. Blocking Selenium calls are delegated to background threads so multiple results can be processed in parallel without freezing the async loop.

Step 4: Evidence Extraction (LLM) The evidence extractor analyzes retrieved content against claims using structured stance definitions:

- **SUPPORTS:** Evidence confirms or validates the claim through direct statements, semantic equivalents, or mechanism descriptions.
- **REFUTES:** Evidence contradicts or disproves the claim through counter-evidence or statements that evidence is unproven/disproven.
- **MIXED:** Both supporting and refuting elements present.
- **UNCLEAR:** Genuinely ambiguous content that discusses related topics without addressing the specific claim.

Critical prompt instructions ensure that mere discussion equals UNCLEAR, that mechanisms are recognized even without exact terminology, and that both Google snippets and page content are considered with better evidence prioritized. Importantly, the evidence extractor pays special attention to qualifiers such as “only”, “never”, “always”, and temporal scope limitations (e.g., “since X date”)—this enables detection of claims that may be technically accurate but misleadingly framed due to omitted context or overgeneralization.

Adaptive Credibility Filtering The search orchestrator implements adaptive credibility filtering as a key innovation for ensuring evidence quality. The algorithm operates in three phases:

1. **Initial Batch:** Fetch $2\times$ the desired limit to provide filtering margin
2. **Quality Check:** If $>50\%$ of results are unreliable, fetch an additional batch to increase the pool of high-quality sources
3. **Intelligent Filtering:** Sort all results by reliability score and select the top reliable sources, with a minimum guarantee ensuring at least some results are returned even if reliability is universally low

Additional filtering mechanisms include stance filtering (removing unclear results if $>50\%$ have definitive stances) and URL deduplication using sets to prevent duplicate sources across different queries for the same claim.

3.3.5 Output Generator

The Output Generator synthesizes evidence into coherent verdicts with confidence levels, quality scoring, and timestamp mapping.

Two-Step Process The Output Generator employs a hybrid approach combining algorithmic evidence organization with LLM-based synthesis:

Step 1: Build Evidence Bundle (Algorithmic) — The system groups evidence by stance (supports, refutes, mixed, unclear), deduplicates sources by URL, and sorts within each group by reliability rating (high first), then numerical score, with alphabetic tie-breaking for consistency.

Step 2: Generate Verdict (LLM) — Organized evidence is formatted into a structured prompt containing stance labels, source counts, reliability ratings, and evidence summaries. The system prompt instructs the LLM to synthesize concise verdicts, naming sources explicitly only when clarifying contrasting perspectives or when evidence directly conflicts. This reduces verbosity while maintaining attribution transparency.

Evidence Quality Score (Algorithmic) The quality score is calculated algorithmically without LLM involvement for consistency and speed. The scoring formula combines three components with different weights: a base score of 0.3 for having any evidence, an actionable stance bonus of up to 0.3 (scaled by the number of supports/refutes/mixed sources, saturating at 3 sources), and a reliability bonus of up to 0.4 (scaled by the number of high/medium reliability sources, saturating at 3 sources). This design prioritizes both actionable stances and source reliability, with the maximum achievable score of 1.0 indicating high-quality, decisive evidence from multiple reliable sources.

Table 5 summarizes the quality score components.

Table 5: Evidence quality score breakdown

Component	Weight	Criteria
Base	0.3	Having any evidence
Actionable	0.3	Up to 3 supports/refutes/mixed sources
Reliability	0.4	Up to 3 high/medium reliability sources
Maximum	1.0	

Integrated Verdict Generation (within Level 1) Verdicts are generated immediately after each claim’s evidence collection completes, within the same parallel execution context as the claim processing. This design choice eliminates the latency overhead of waiting for all claims to finish evidence collection before beginning verdict synthesis. Each claim’s verdict generation executes as soon as its evidence is ready, allowing early-finishing claims to produce results while slower claims continue processing. After all parallel claim tasks complete, the reports are sorted by evidence quality score (descending) to prioritize high-confidence verdicts in the user interface.

3.4 Latency Optimization Strategies

The pipeline ships with a concrete set of latency optimizations. The following subsections describe what was actually built: tuned model choices, strict token budgets, trimmed inputs, single-call components, parallel execution, real-time streaming updates, and classical fallbacks for non-reasoning tasks. Together these measures reduced mean end-to-end latency to 129.6 seconds for the evaluation set.

3.4.1 Model Selection

The default model (gpt-4o-mini) is selected for its balanced performance across speed, cost-effectiveness, and large context window (128K tokens). This model provides sufficient reasoning capabilities for fact-checking tasks while maintaining low latency and competitive pricing (\$0.15 per million input tokens, \$0.60 per million output tokens). For budget-conscious deployments or offline operation, local Ollama models (qwen3:8b, qwen3:4b) offer zero-cost inference at the expense of potential quality degradation.

3.4.2 Output Constraints

Each component has carefully tuned `max_tokens` limits to minimize generation latency and API costs without sacrificing information quality. Claims are limited to approximately 40 words (sufficient for most factual assertions), context descriptions to 20 words (brief background), and evidence summaries to 1–2 sentences (key findings only). These constraints are enforced through explicit prompt instructions and validated against output token limits.

3.4.3 Content Trimming

Input token counts are minimized through aggressive content trimming strategies. Web content is trimmed to 6,000–8,000 characters before being passed to the Evidence Extractor, removing excessive context while retaining the most relevant portions (typically the first several paragraphs of an article). Evidence prompts include only Google snippets and extracted page text, explicitly excluding raw HTML, JavaScript, CSS, and other non-content elements that would inflate token counts without improving extraction quality.

3.4.4 Combined Operations

The pipeline minimizes LLM API calls by combining operations into single requests wherever possible. Each component makes exactly one LLM call per input unit (one call for claim extraction, one per claim for query generation, one per search result for evidence extraction, and one per claim for verdict synthesis), with no multi-turn conversations that would multiply request counts.

3.4.5 Three-Level Async Architecture

The system implements three levels of nested parallelization to maximize throughput while maintaining dependency ordering:

- **Level 1 (Claims):** After extracting N claims from the transcript, all claims are processed in parallel. Each claim independently proceeds through query generation, evidence search, and verdict generation. Critically, each claim’s verdict is generated immediately after its evidence collection completes, rather than waiting for all claims to finish—this optimization reduces perceived latency by producing results progressively.
- **Level 2 (Queries per Claim):** Within each claim’s processing, the Query Generator produces M queries. These queries are executed in parallel, enabling simultaneous search across different query formulations (direct, alternative, source-seeking, contextual).

- **Level 3 (Search Results per Query):** Within each query’s execution, the Online Search component retrieves K results from Google. The four-step pipeline (reliability assessment, content fetching, and evidence extraction) runs in parallel for all K results, with each result processed independently.

This design achieves maximum theoretical parallelization of $N \times M \times K$ operations during the search phase, bounded only by system resources and API rate limits.

3.4.6 Real-Time Streaming

Server-Sent Events (SSE) provide progressive updates as the pipeline executes, improving perceived responsiveness for users. Each update labels the active stage (e.g., “claim extraction”, “processing_claim_2”, “generating_report”) alongside a percentage so the UI can display the exact component currently running. Extracted claims are streamed immediately after the Claim Extractor finishes, enabling preview and early user feedback while later stages gather evidence.

3.4.7 Classical Methods for Non-Reasoning Tasks

Table 6 shows operations handled by classical algorithms rather than LLMs.

Table 6: Operations using classical methods

Operation	Method	Rationale
Reliability scoring	Rule-based + MBFC lookup	Faster, deterministic, no API cost
Claim localization	Fuzzy string matching	No LLM needed for text search
Evidence quality score	Algorithmic calculation	Consistent, fast, reproducible
URL deduplication	Hash set	$O(1)$ lookup
Stance filtering	Threshold-based	Simple percentage check

3.5 Experimentation and Evaluation Framework

Evaluating LLM-based fact-checking systems presents unique challenges: outputs are non-deterministic, external dependencies (web search) introduce variability, and traditional benchmarks risk overfitting [28]. To address these challenges, a three-component experimentation framework was developed to capture complete execution traces, support batch experimentation, and compute performance metrics.

3.5.1 Framework Architecture

The framework follows a linear data flow through three stages:

1. **Experiment Runner:** Executes the fact-checking pipeline on configured videos, invoking the tracking module for each run.
2. **Tracking Module:** Captures all execution data—inputs, outputs, LLM calls, timing, and costs—saving structured artifacts for later analysis.
3. **Evaluator:** Computes performance metrics including comparisons against ground truth annotations, system efficiency measurements, and source quality assessments.

3.5.2 Tracking Module

The tracking module implements a singleton pattern with context manager support, enabling any pipeline component to log data without explicit parameter passing. When initialized, the tracker creates a timestamped run directory and registers itself as the global tracker. The context manager protocol ensures automatic saving on exit.

Each run generates structured artifacts containing: run configuration and parameters, complete records of all LLM calls with prompts, responses, latency and cost, final extracted claims and fact-check verdicts, aggregated timing and cost metrics, and the original video transcript for reference.

LLM call tracking is achieved via a decorator that instruments the Pydantic AI agent methods, automatically recording component name, model, timestamp, latency, token counts, and calculated cost for every inference call.

3.5.3 Evaluator

The evaluator computes performance metrics using modular components for each evaluation dimension:

- **Claim extraction metrics:** Precision@k, Recall, F1, and MAP computed using semantic similarity matching between extracted and ground truth claims.
- **Verdict accuracy:** Comparison of system stances against ground truth labels.
- **Evidence retrieval metrics:** Success rate and source reliability distribution based on MBFC credibility ratings.
- **System efficiency:** Latency and cost aggregation across all pipeline components.

- **LLM-as-judge evaluation** (optional): LLM-as-judge is an evaluation paradigm where a language model assesses the quality of outputs from another model (or the same model), providing scores or judgments on dimensions difficult to capture with traditional metrics [27]. In this system, LLM-as-judge can evaluate claim relevance, evidence quality, and verdict coherence.

Evaluations execute in parallel across multiple videos for efficiency. Results include per-video reports and aggregate statistics (means, standard deviations, distributions) across all evaluated videos.

3.6 API Layer and User Interface

3.6.1 FastAPI Setup

The API uses FastAPI with CORS middleware configured for local frontend development, supporting common development server ports. API routes are organized under a versioned prefix to enable future API evolution without breaking existing clients.

3.6.2 Streaming Endpoint with SSE

The streaming endpoint exposes the fact-checking run as a server-sent event stream. Progress callbacks enqueue updates that the API emits in real time, allowing clients to maintain a single open connection while receiving incremental status messages. Events follow a fixed sequence (transcript extraction, claim extraction, per-claim processing, verdict generation, completion) so the frontend can map each update to user-facing milestones without polling.

3.6.3 User Interface

A web-based frontend provides an accessible interface for end users to interact with the fact-checking pipeline. The interface is built with React and communicates with the backend via the streaming API endpoint. Figure 3 presents the four main interface states.




Figure 3: User interface screenshots showing the fact-checking workflow: (a) users paste a YouTube URL, (b) real-time progress is displayed during analysis, (c) results are presented alongside the embedded video, and (d) each claim shows its verdict, confidence level, and supporting evidence with source reliability ratings.

The interface shows transparency by displaying source reliability ratings, confidence levels, and direct links to evidence sources. Users can click “Jump” buttons to navigate directly to the video timestamp where each claim was made, enabling quick verification of the original context.

3.7 Engineering Practices

The thesis emphasizes data science outcomes, yet maintaining consistent engineering habits keeps experiments reproducible and easier to debug. Practices are kept lightweight and focused on what directly supports the fact-checking pipeline:

- **Typed schemas:** Pydantic models plus strict mypy checks prevent interface drift between components while supplying clear validation errors during runs.
- **Logging & fallbacks:** Module-level logging records claim/query context, and simple retry rules (e.g., fall back to Google snippets when scraping fails) keep long runs from

collapsing on transient issues.

- **Centralized configuration:** Environment variables store secrets, YAML files define experiment batches, and a small Python settings module captures model limits, making it easy to reproduce or tweak runs.
- **Async + background threads:** Claims, queries, and search results fan out via asyncio gather calls, while Selenium/WHOIS work shifts to threads so blocking I/O never stalls the event loop; the same progress callbacks drive the SSE stream for user feedback.
- **Pre-commit hooks:** Lightweight automated checks run before every commit to keep formatting, linting, and type rules consistent.

4 Results

This section presents the experimental evaluation of Factible across 30 YouTube videos spanning diverse topics including health, science, politics, and climate. The evaluation framework follows established practices from claim detection and fact-checking research [9, 32], combining ground truth comparison with LLM-as-judge quality assessments. Every run limited the pipeline to five claims per video, a setting chosen after analyzing the precision-recall tradeoff across multiple configurations while balancing cost and latency constraints.

4.1 Experimental Setup

4.1.1 Evaluation Dataset

The evaluation corpus consists of 30 YouTube videos manually annotated with ground truth claims. This sample size is comparable to evaluation scales used in end-to-end fact-checking systems; for example, ClaimBuster evaluated their system on 25 presidential debates [?]. Videos were selected across three thematic categories—climate, health, and political/social issues—with 10 videos per category to ensure balanced representation. Within each category, videos were equally split between factual content (5 videos) and misinformation (5 videos), allowing evaluation of the system’s performance across different truth orientations. Videos range from educational science content to political commentary, representing diverse domains and claim densities. Table 7 summarizes the dataset characteristics.

Table 7: Evaluation dataset statistics

Metric	Value
Total videos	30
Ground truth claims per video (mean)	16.8
Ground truth claims per video (range)	9–35
Total ground truth claims	503
System claims extracted per video	5

4.1.2 Ground Truth Annotation

Ground truth annotations were created following a structured protocol. For each video, all factual, verifiable claims from the transcript were manually annotated. Each claim was annotated with an importance score (0.0–1.0) reflecting its centrality to the video’s main argument, and a verdict label indicating the expected verification outcome (SUPPORTS, REFUTES, MIXED, or UNCLEAR). The annotation process followed guidelines from ClaimBuster’s check-worthiness criteria [9], prioritizing claims that are specific, verifiable, and consequential. The complete evaluation dataset, including all 30 annotated videos with their ground truth claims, importance ratings, and expected verdicts, is provided in Appendix A.

4.1.3 Evaluation Metrics

The evaluation employs metrics at two levels: claim extraction quality and verdict accuracy.

Claim Alignment via Semantic Similarity System claims are matched to ground truth using sentence-transformer embeddings (all-MiniLM-L6-v2). Claims whose cosine similarity exceeds 0.7 count as matches; others become false positives or false negatives. A greedy pass ensures each ground truth claim pairs with at most one system claim. This simple semantic alignment yields the true positives needed for precision, recall, F1, and MAP calculations.

Claim Extraction Metrics:

- **Precision@k:** Proportion of extracted claims matching any ground truth claim, using semantic similarity matching with a threshold of 0.7. This metric follows the standard information retrieval formulation used in claim detection systems [11].
- **Recall@k:** Proportion of ground truth claims matched by the top- k extracted claims [11].
- **F1 Score:** Harmonic mean of precision and recall.

- **Mean Average Precision (MAP):** Ranking quality metric from information retrieval, measuring whether important claims appear early in the extraction order [11].
- **Recall@Important:** Recall specifically for high-importance claims ($\text{importance} \geq 0.80$).
- **Importance-Weighted Coverage:** Percentage of total ground truth importance mass captured by matched claims.

Verdict Accuracy Metrics:

- **Stance Accuracy:** Classification accuracy for the four-class stance problem (SUPPORTS, REFUTES, MIXED, UNCLEAR), measuring the percentage of verdicts matching ground truth stance.

System Efficiency Metrics:

- **Latency:** End-to-end processing time per video.
- **Evidence Retrieval Rate:** Proportion of queries successfully retrieving evidence.
- **Source Reliability:** Distribution of source reliability ratings across retrieved evidence.

4.1.4 Component Evaluation Scope

The evaluation focuses on components where the system makes reasoning decisions that can be compared against ground truth. Specifically, the evaluation covers **Claim Extraction** (precision, recall, importance ranking) and **Verdict Generation** (stance accuracy, explanation quality) as these components employ LLM-based reasoning that can produce varying results.

Components relying on external APIs—**Transcript Extraction** (YouTube Transcript API) and **Online Search** (Serper/Google)—are not evaluated directly, as their performance depends on third-party services rather than the system’s design. However, their effectiveness is implicitly covered by the end-to-end evaluation: if transcript extraction fails, no claims can be extracted; if search fails, verdict accuracy degrades. The **Query Generator** component is assessed indirectly through evidence retrieval success rates, as effective queries should yield relevant evidence.


4.2 Precision-Recall Tradeoff Analysis

Before presenting detailed results, this section analyzes the precision-recall tradeoff to justify the choice of the configuration of maximum claims to fetch to be 5 (`max_claims=5`) as the primary configuration. The system was evaluated across six configurations with `max_claims` $\in \{1, 3, 5, 7, 10, 15\}$.

Table 8: Precision-recall tradeoff across different `max_claims` configurations

<code>max_claims</code>	Precision	Recall	F1	MAP
1	0.800	0.052	0.098	0.800
3	0.822	0.159	0.264	0.897
5	0.813	0.262	0.390	0.870
7	0.795	0.359	0.486	0.862
10	0.769	0.471	0.573	0.854
15	0.719	0.570	0.623	0.865

Figure 4 illustrates the precision-recall tradeoff. As `max_claims` increases, recall improves from 5.2% to 57.0%, while precision decreases from 82.2% to 71.9%. The slight precision increase from $k = 1$ (80.0%) to $k = 3$ (82.2%) indicates that the system benefits from extracting multiple high-confidence claims rather than being forced to select exactly one. Beyond $k = 5$, the classic precision-recall tradeoff becomes pronounced.

Figure 4: Precision-recall curve for different `max_claims` values. The selected operating point ($k = 5$) achieves 81.3% precision at 26.2% recall, balancing claim quality with coverage.

The configuration `max_claims=5` was selected as the primary operating point because it maintains high precision (81.3%) while achieving reasonable recall (26.2%), achieves a strong MAP score (0.870) indicating good ranking quality, and provides favorable cost and latency

characteristics. Beyond `max_claims=5`, processing time and API costs increase linearly with claim count, while precision degrades. For a fact-checking application where user trust depends on accuracy and responsiveness, balancing precision, cost, and latency is critical—presenting 5 high-quality claims efficiently is more valuable than presenting 15 claims with higher false positive rates, increased costs, and longer wait times.

4.3 Claim Extraction Performance

Table 9 presents the claim extraction results at the primary configuration (`max_claims=5`).

Table 9: Claim extraction performance (n=30 videos, `max_claims=5`)

Metric	Mean	Std Dev
Precision@5	0.813	0.171
Recall	0.262	0.088
F1 Score	0.390	0.111
Mean Average Precision (MAP)	0.870	0.146
Recall@Important (≥ 0.80)	0.395	0.178
Importance-Weighted Coverage	0.292	0.096
Importance MAE	0.126	0.060

4.3.1 Interpretation of Results

The system achieves 81.3% precision, meaning that approximately 4 out of 5 extracted claims on average match ground truth claims. This high precision indicates that the claim extractor successfully identifies legitimate factual claims rather than extracting irrelevant or fabricated statements.

The overall recall of 26.2% reflects the constraint of extracting only 5 claims from videos averaging 16.8 ground truth claims. However, the Recall@Important metric (39.5%) demonstrates that the system prioritizes high-importance claims—capturing nearly 40% of the most critical claims while only extracting approximately 30% of the total claim set. This selective extraction behavior aligns with the design goal of thesis-first reasoning, where claims are ranked by their impact on the video’s central argument.

The MAP score of 0.870 indicates strong ranking quality: important claims consistently appear early in the extraction order. This metric, borrowed from ClaimBuster’s evaluation framework [9], validates that the importance scoring mechanism effectively prioritizes claims.

The importance MAE of 0.126 (on a 0–1 scale) shows that system-assigned importance scores

closely approximate human judgments, with typical errors of approximately one importance tier (e.g., scoring a claim 0.7 when ground truth is 0.85).

4.3.2 Contextualizing Recall

The 26.2% recall, while appearing low in isolation, must be interpreted in context. Given that videos contain an average of 16.8 checkable claims and the system extracts 5, a theoretical maximum recall of approximately 30% exists under this constraint. The achieved recall of 26.2% therefore represents strong performance relative to the configuration limit.

Furthermore, the importance-weighted coverage of 29.2% indicates that the system captures nearly one-third of the total “importance mass” of ground truth claims. For practical fact-checking applications where user attention is limited, presenting 5 high-quality, important claims provides more value than exhaustive but overwhelming coverage.

4.4 Verdict Generation Performance

Table 10 presents the verdict accuracy results.

Table 10: Verdict generation performance (n=30 videos)

Metric	Mean	Std Dev
Stance Accuracy	0.733	0.334

4.4.1 Accuracy Distribution Analysis

The standard deviation (0.334) in verdict accuracy warrants investigation. Analysis of per-video results reveals a distribution skewed toward high accuracy:

- **21 videos (70.0%)** achieved high accuracy ($\geq 75\%$)—the majority of extracted claims were correctly classified.
- **5 videos (16.7%)** achieved medium accuracy (25–74%)—partial verdict correctness.
- **4 videos (13.3%)** achieved low accuracy ($< 25\%$)—most claims were incorrectly classified.

Evidence Retrieval and Verdict Quality With an evidence retrieval success rate of 94.7%, the system consistently finds relevant sources for most claims. Table 11 illustrates how verdict accuracy varies across the dataset.

Table 11: Verdict accuracy distribution with evidence retrieval rates

Video Topic	Retrieval Rate	Verdict Acc.	Avg Sources
<i>High accuracy ($\geq 75\%$) — 21 videos</i>			
Brain Benefits of Exercise (TED)	100%	100%	2.5
Climate Change (Nate Geo)	100%	100%	2.3
Fossil Fuels	100%	100%	2.7
UK Election Results	100%	75%	2.8
<i>Medium/Low accuracy ($< 75\%$) — 9 videos</i>			
Gender Wage Gap	100%	50%	2.7
FBI & January 6th	100%	50%	2.3
Inflation Explainer	60%	25%	1.8
Immigration Statistics	100%	0%	2.9

Notably, successful evidence retrieval does not guarantee high verdict accuracy. Some politically contentious topics achieve 100% retrieval but lower verdict accuracy, suggesting that the challenge lies not in finding evidence but in correctly synthesizing conflicting sources or matching the ground truth annotator’s interpretation.

Factors Affecting Verdict Accuracy Analysis of low-accuracy videos reveals several contributing factors:

- **Contested claims:** Topics with legitimate disagreement (e.g., wage gap interpretations, immigration statistics) may have evidence supporting multiple stances, making definitive verdicts challenging.
- **Ground truth subjectivity:** Some claims involve nuanced interpretations where reasonable annotators might disagree on the correct stance.
- **Evidence-claim mismatch:** Retrieved evidence may address related but not identical claims, leading to verdict errors.

These findings suggest that further improvements require enhanced evidence synthesis and more sophisticated handling of contested claims, rather than simply improving retrieval success.

4.4.2 Comparison to Random Baseline

The stance accuracy of 73.3% substantially exceeds a random baseline. For a four-class classification problem (SUPPORTS, REFUTES, MIXED, UNCLEAR), random guessing would achieve approximately 25% accuracy. The system’s 73.3% accuracy represents a $2.93\times$ improvement over random, demonstrating meaningful verification capability. Moreover, unlike random classification, the system provides evidence-backed explanations that enable users to evaluate the verdict’s reasoning.

4.5 Evidence Retrieval Performance

4.5.1 Retrieval Success

Table 12 summarizes evidence retrieval performance.

Table 12: Evidence retrieval performance (n=30 videos)

Metric	Value
Evidence retrieval success rate	94.7%
Average sources per query	1.52
Average evidence items per claim	2.41

The evidence retrieval success rate of 94.7% indicates that the vast majority of search queries return usable evidence. When evidence is retrieved, claims receive an average of 2.41 evidence items, providing multiple perspectives for verdict synthesis.

4.5.2 Source Reliability Distribution

A critical aspect of fact-checking is source quality. Table 13 presents the distribution of source reliability ratings across all retrieved evidence.

Table 13: Source reliability distribution (n=724 total sources)

Reliability Rating	Count	Percentage
High	605	83.6%
Medium	110	15.2%
Low	0	0.0%
Unknown	9	1.2%

The overwhelming majority of retrieved sources (83.6%) receive high reliability ratings from the Media Bias/Fact Check-based assessment system [21]. No sources received low reliability ratings, and only 1.2% were classified as unknown (typically due to missing MBFC data for niche domains). This distribution reflects both the reliability scoring heuristics and the retry behavior that fetches additional results whenever the initial batch skews toward low-reliability domains.

4.6 System Efficiency

4.6.1 Processing Latency

Table 14 presents processing time statistics.

Table 14: System latency (n=30 videos)

Metric	Value
Mean latency	129.6 seconds
Standard deviation	101.8 seconds
Minimum latency	36.5 seconds
Maximum latency	643.7 seconds
Total processing time (30 videos)	64.8 minutes

The mean processing time of 129.6 seconds per video enables near-interactive use for individual videos. The high variance (standard deviation 101.8s) reflects multiple contributing factors:

- **Video length:** Longer transcripts require more LLM tokens for claim extraction, increasing inference time.
- **Pipeline parameters:** The configuration parameters `max_claims`, `max_queries`, and `max_results_per_query` directly multiply the number of downstream operations. With the evaluation configuration (`max_claims = 5`, `max_queries = 3`, `max_results = 3`), each video triggers up to $5 \times 3 \times 3 = 45$ evidence extraction operations.
- **Evidence retrieval success:** Videos with failed evidence retrieval complete faster (fewer web requests), while videos requiring multiple successful searches experience longer latencies.
- **Web scraping variability:** JavaScript-heavy sites require longer Selenium wait times, and some domains respond slower than others.

4.6.2 Cost Analysis

All experiments used GPT-4o-mini for LLM inference at current pricing (\$0.15 per million input tokens, \$0.60 per million output tokens). Across the 30-video evaluation corpus, the average cost per video was \$0.003, with individual videos ranging from \$0.0008 to \$0.0164 depending on transcript length, claim complexity, and evidence retrieval needs. The total cost for processing all 30 videos was \$0.09, demonstrating the practical affordability of the system for individual users. This cost-effectiveness contrasts with concerns about LLM deployment costs noted in prior work [?], showing that fact-checking systems can achieve meaningful accuracy at minimal expense when using appropriately-sized models.

4.7 Qualitative Analysis

4.7.1 Successful Extraction Examples

Table 15 presents examples of successful claim extractions demonstrating semantic matching between ground truth and system-extracted claims.

Table 15: Examples of successful claim extraction with semantic matching

Ground Truth Claim	System-Extracted Claim	Imp.
A single workout immediately increases levels of neurotransmitters like dopamine, serotonin, and norepinephrine	A single workout increases levels of neurotransmitters like dopamine, serotonin, and norepinephrine	0.9
HIV infects one of the immune cells that is central to the body's response to pathogens—the helper T-cell	HIV infects helper T-cells, which are central to the immune response	0.95
Scientists at UCT have uncovered garlic's cancer fighting properties	Scientists at UCT uncovered garlic's cancer-fighting properties	0.95

These examples demonstrate that the system successfully extracts claims while allowing minor paraphrasing and condensation. The semantic similarity matching correctly identifies these as equivalent claims despite surface-level textual differences.

4.7.2 Error Analysis: Verdict Failures

Analysis of verdict errors reveals systematic patterns. With high evidence retrieval success (94.7%), the primary failure modes involve stance misclassification and handling nuanced claims where conflicting evidence requires domain expertise to synthesize correctly. Table 16 categorizes the primary error types.

Table 16: Verdict error categories

Error Type	Description
Evidence retrieval failure	No evidence retrieved; system defaults to UNCLEAR
Stance misclassification	Evidence retrieved but stance incorrectly assessed (e.g., MIXED classified as REFUTES)
Nuanced claims	Claims requiring domain expertise to evaluate mixed evidence

Error analysis reveals that evidence retrieval success does not guarantee verdict accuracy. Two of the four lowest-accuracy videos achieved 100% evidence retrieval but 0% verdict accuracy, indicating that the challenge lies in evidence synthesis rather than retrieval. These cases involve politically contested claims where ground truth stances require nuanced interpretation of conflicting sources.

4.7.3 Analysis of Low-Accuracy Cases

Only 4 videos (13.3%) achieved less than 25% verdict accuracy. Detailed analysis reveals their characteristics:

Table 17: Low-accuracy video analysis (<25% verdict accuracy)

Video Topic	Retrieval	Accuracy	Likely Cause
Trump's National Emergency	60%	0%	Political claims with contested interpretations
Immigration Statistics	100%	0%	Conflicting sources on contested statistics
Juice vs. Whole Fruit	20%	20%	Limited evidence retrieval
Sleep & Teenage Brain	100%	20%	Nuanced scientific claims

Key observations:

- **High retrieval does not guarantee accuracy:** Two videos achieved 100% retrieval but 0–20% accuracy, confirming that evidence synthesis is the bottleneck for contested claims.
- **No single category dominates:** Low-accuracy videos span both political (2) and health (2) topics, and include both factual content (3) and misinformation (1).
- **Contested claims are hardest:** The common thread is claims where reasonable sources disagree or where ground truth requires nuanced interpretation.

Successful categories: Videos on scientific explanations (e.g., climate science, exercise benefits), health misinformation debunking (e.g., fluoride claims, detox myths), and conspiracy content (e.g., chemtrails, geoengineering) achieved high accuracy, demonstrating the system’s effectiveness when evidence clearly supports or refutes claims.

4.8 Considerations for Generative AI Systems

It is important to contextualize these results within the unique characteristics of generative AI systems. Unlike traditional machine learning models with deterministic outputs, LLM-based systems introduce inherent variability that affects evaluation interpretation.

4.8.1 Non-Determinism in LLM Systems

Despite configuring all LLM calls with `temperature=0.0` to minimize output variability, complete determinism is not guaranteed. Even with zero temperature, LLM outputs may vary across runs due to:

- **Floating-point precision:** GPU computation introduces subtle numerical variations that can cascade through token selection.
- **Model updates:** Cloud-hosted models (e.g., GPT-4o-mini) may be silently updated by providers, affecting outputs over time.
- **Batching effects:** Different batch sizes or concurrent requests may influence internal state.

This inherent non-determinism means that exact reproduction of results is challenging, though setting temperature to zero substantially reduces variability compared to default settings.

4.8.2 External Dependencies and Temporal Sensitivity

Beyond LLM variability, the system's reliance on external web search introduces additional sources of result variability:

- **Search result volatility:** Web search results change over time as new content is indexed and rankings evolve.
- **Content availability:** Websites may become unavailable, paywalled, or block automated access.
- **Rate limiting:** Search APIs may throttle requests, causing some queries to fail during high-load evaluation runs.

These factors contribute to verdict accuracy variance: search results may differ between runs, and a claim that retrieved limited evidence in one execution might find more sources in another.

4.8.3 Implications for Metric Interpretation

Unlike traditional classification tasks where metrics are stable given fixed test data, LLM-based systems produce metrics with inherent variance. When evaluating generative AI systems, researchers should consider:

- **Expected variability:** Metrics may vary by several percentage points across identical evaluation runs.
- **Qualitative validation:** Beyond aggregate metrics, examining individual outputs provides crucial insight into system behavior.

- **Temporal context:** Results reflect system performance at a specific point in time with then-current search results and model versions.

The strategies employed in this work to maximize reproducibility—temperature=0.0 for all LLM calls, fixed random seeds, comprehensive logging, and experiment versioning—represent current best practices for LLM evaluation but cannot eliminate all sources of variability.

4.9 Summary of Key Findings

The experimental evaluation demonstrates that Factible achieves reliable fact-checking performance across diverse video content:

1. **High-precision claim extraction:** 81.3% precision with strong importance ranking (MAP 0.870), meaning users can trust that presented claims are legitimate, high-priority factual statements.
2. **Robust evidence retrieval:** 94.7% success rate with 83.6% of sources from high-reliability origins, demonstrating effective query generation and source filtering.
3. **Consistent verdict accuracy:** 73.3% overall accuracy, with 70% of videos (21/30) achieving $\geq 75\%$ accuracy. This represents a $2.93\times$ improvement over random baselines.
4. **Identified challenges:** The 4 low-accuracy videos (13.3%) involve politically contested claims with legitimate disagreement, where evidence synthesis rather than retrieval poses the challenge.
5. **Practical efficiency:** 129.6 seconds mean latency at \$0.003 per video enables cost-effective, near-interactive use.

These results position Factible as a reliable tool for preliminary fact-checking of YouTube content. The system performs well on scientific, health, and clearly verifiable claims, while contested political topics with conflicting evidence sources represent the primary remaining challenge for future work.

5 Conclusions and Future Work

This section presents the conclusions of this thesis, evaluates the achievement of initial objectives, reflects on the methodology employed, discusses ethical and sustainability considerations, and proposes directions for future research.

5.1 Summary of Contributions

This thesis has presented Factible, a multi-agent system for automated fact-checking of YouTube videos. The work makes several contributions to the field of automated fact-checking and demonstrates the practical viability of LLM-based verification systems.

5.1.1 Technical Contributions

1. **End-to-End Video Fact-Checking Pipeline:** Unlike prior research that focuses on isolated components, Factible implements a complete pipeline from YouTube URL to verified claims with evidence-backed verdicts. The system handles transcript extraction, claim detection, query generation, evidence retrieval, and verdict synthesis autonomously.
2. **Thesis-First Claim Extraction:** A novel prompting approach that prioritizes claims based on their impact on the video's central argument, achieving 81.3% precision and 0.870 MAP score for claim extraction.
3. **Open-Web Evidence Retrieval with Credibility Assessment:** The system retrieves evidence from the live web rather than closed knowledge bases, incorporating source credibility evaluation using Media Bias/Fact Check data and domain heuristics. This enables verification of claims about current events and specialized topics.
4. **Three-Level Parallel Architecture:** An efficient execution model that parallelizes processing across claims, queries, and search results, achieving mean processing time of 129.6 seconds per video while maintaining cost under \$0.01 per video.
5. **Production-Ready Implementation:** The complete system includes a REST API with SSE streaming, a React-based web interface, and comprehensive experimentation infrastructure, bridging the gap between research prototypes and practical tools.

5.1.2 Research Contributions

1. **Annotated Evaluation Dataset:** A corpus of 30 YouTube videos with 503 manually annotated ground truth claims across health, climate, and political topics, providing a foundation for future research.
2. **Evaluation Framework:** A comprehensive methodology combining precision/recall metrics, semantic similarity matching, verdict accuracy assessment, and system efficiency measurements appropriate for LLM-based fact-checking evaluation.

3. **Analysis of Failure Modes:** Systematic identification of where automated fact-checking struggles, particularly with contested political claims where evidence synthesis rather than retrieval poses the challenge.

5.2 Achievement of Objectives

This section evaluates the degree to which the initial objectives defined in Section 1 were achieved.

5.2.1 Main Objective

Objective: Design, implement, and evaluate a multi-agent system capable of automatically verifying factual claims in YouTube videos through retrieval and analysis of web-based evidence, generating substantiated verdicts with confidence estimates.

Achievement: Fully achieved. The system processes YouTube videos end-to-end, extracting claims, retrieving evidence from multiple web sources, and generating verdicts with confidence levels. Evaluation on 30 videos demonstrates 81.3% claim extraction precision, 94.7% evidence retrieval success, and 73.3% verdict accuracy.

5.2.2 Specific Objectives

1. **Multi-agent architecture design: Achieved.** Five specialized components with well-defined Pydantic schemas operate through a modular pipeline, enabling independent testing and optimization.
2. **System component implementation: Achieved.** All five components (Transcriptor, Claim Extractor, Query Generator, Online Search, Output Generator) are fully implemented and operational.
3. **LLM usage optimization: Partially achieved.** The system uses GPT-4o-mini effectively with deterministic settings and token budgets. Full comparison with open-source models (Qwen, DeepSeek) was not completed due to scope constraints, though the infrastructure supports such comparisons.
4. **System evaluation: Achieved.** Comprehensive evaluation including precision/recall metrics, verdict accuracy, evidence retrieval statistics, latency, cost analysis, and qualitative case studies.
5. **End-to-end system implementation: Achieved.** REST API with SSE streaming, React web interface, and documentation are complete. Source code is publicly available.

6. **Ethical and bias considerations: Achieved.** The system incorporates source transparency, multi-perspective evidence presentation, and clear confidence levels. Discussion of limitations and bias mitigation strategies is included.

5.3 Critical Assessment of Methodology

5.3.1 Strengths of the Approach

The Design and Creation methodology [12, 25] proved appropriate for this project:

- **Iterative development** enabled rapid identification and resolution of component issues, particularly in prompt engineering and evidence extraction.
- **Modular architecture** facilitated independent testing and optimization of each pipeline component.
- **Practical focus** ensured the system addresses real-world requirements (latency, cost, usability) beyond academic benchmarks.
- **Comprehensive evaluation** combining quantitative metrics with qualitative analysis provided meaningful insights into system behavior.

5.3.2 Limitations of the Approach

Several limitations should be acknowledged:

- **Single annotator for ground truth:** All 503 ground truth claims were annotated by a single person (the author), introducing potential bias. Inter-annotator agreement studies would strengthen the evaluation.
- **Dataset size:** While 30 videos with 503 claims is comparable to prior work, a larger dataset would enable more robust statistical conclusions and better representation of edge cases.
- **English-only scope:** The system currently supports only English content, limiting applicability in multilingual contexts where misinformation is equally prevalent.
- **Temporal snapshot:** Evaluation was conducted at a specific point in time (December 2025). Both web search results and LLM model versions change, affecting reproducibility.
- **Limited model comparison:** Full comparison between commercial and open-source LLMs was not completed, leaving optimization opportunities unexplored.

5.4 Reflection on Ethical and Sustainability Considerations

5.4.1 Sustainability Assessment

The project's environmental impact was considered throughout development:

- **Computational efficiency:** Using GPT-4o-mini rather than larger models (GPT-4, GPT-4-turbo) significantly reduces energy consumption while maintaining adequate quality. The average cost of \$0.003 per video reflects both financial and environmental efficiency.
- **Optimization strategies:** Token budgets, content trimming, and classical algorithms for non-reasoning tasks minimize unnecessary LLM calls.
- **Local model support:** The system supports Ollama-based local models, enabling zero-cloud-cost operation for users with appropriate hardware, further reducing the system's carbon footprint for some deployments.

While LLM usage inherently involves energy consumption, the system's design choices represent a conscious effort to balance capability with environmental responsibility.

5.4.2 Ethical Considerations

The system addresses ethical concerns through several mechanisms:

- **Transparency:** All evidence sources are exposed with URLs, reliability ratings, and stance classifications, enabling users to verify the system's reasoning.
- **Multi-perspective presentation:** Evidence is organized by stance (supports, refutes, mixed, unclear), avoiding false certainty and acknowledging when claims are genuinely contested.
- **Confidence levels:** Verdicts include confidence ratings (low, medium, high), helping users calibrate trust appropriately.
- **Tool positioning:** The system is positioned as a fact-checking assistant rather than an authoritative source, emphasizing its role in supporting rather than replacing human judgment.

5.4.3 Potential Risks and Mitigation

Several risks warrant ongoing attention:

- **Over-reliance on automation:** Users might accept system verdicts uncritically. Mitigation: Clear messaging that the system provides preliminary assessment, not definitive truth.
- **Bias amplification:** LLMs may perpetuate biases from training data. Mitigation: Source diversity in evidence retrieval and transparent presentation of multiple perspectives.
- **Weaponization concerns:** The system could theoretically be misused to generate plausible-sounding refutations of true claims. Mitigation: Open-source release enables scrutiny, and the evidence-based approach makes manipulation detectable.

5.4.4 Sustainable Development Goals Impact

The project contributes positively to three SDGs:

- **SDG 4 (Quality Education):** The system serves as a media literacy tool, helping users develop critical thinking skills by exposing the fact-checking process.
- **SDG 16 (Peace, Justice and Strong Institutions):** By providing transparent verification mechanisms, the system supports informed public discourse and combats misinformation that erodes institutional trust.
- **SDG 10 (Reduced Inequalities):** The open-source, low-cost design democratizes access to fact-checking capabilities previously available only to well-resourced organizations.

5.5 Lessons Learned

Several key lessons emerged from this project:

1. **Prompt engineering is critical but brittle:** Small changes in prompt wording can significantly affect LLM outputs. The thesis-first approach required extensive iteration to achieve stable, high-quality claim extraction.
2. **Evidence retrieval is easier than evidence synthesis:** The system achieves 94.7% evidence retrieval success but only 73.3% verdict accuracy. The challenge lies in correctly interpreting and synthesizing conflicting sources, not in finding them.

3. **Contested claims remain difficult:** Political and socially contentious topics with legitimate disagreement present fundamental challenges for automated fact-checking that may require different approaches than straightforward factual verification.
4. **Practical constraints matter:** Latency, cost, and usability considerations significantly shaped design decisions. Academic metrics alone are insufficient for evaluating systems intended for real-world use.
5. **Modular design pays dividends:** The five-component architecture enabled independent development, testing, and optimization, proving essential for managing complexity.

5.6 Future Work

While the current implementation demonstrates the viability of automated fact-checking for YouTube videos, several directions warrant further investigation to enhance the system’s capabilities and broaden its applicability.

5.6.1 Multilingual Support

Future iterations should extend transcript extraction, claim detection, and verdict synthesis to major languages beyond English (e.g., Spanish, Portuguese, Hindi, Arabic). Misinformation is a global phenomenon, and the regions with the highest volumes of misinformation often have limited access to fact-checking resources. Key challenges include adapting prompts for different languages, identifying reliable sources in multiple languages, and evaluating source credibility across cultural contexts.

5.6.2 Large Language Model Comparison

Future studies should benchmark multiple commercial and local models (e.g., GPT-4 variants, Claude, Gemini, LLaMA, Qwen, Mistral) across pipeline components. Such comparisons would map the cost, latency, and accuracy trade-offs for mixed deployments where different models handle different components based on their strengths.

5.6.3 Enhanced Prompt Engineering

Maintaining a small prompt library with documented variants, automated prompt tuning, and curated examples per claim type would make it easier to evolve instructions without rewriting code. Techniques like constitutional AI or self-refinement could improve consistency and reduce brittleness.

5.6.4 Enhanced Source Reliability Assessment

Source reliability could be refined by incorporating:

- Publisher-level signals (impact factors, retraction history for academic sources)
- Topic-aware weighting (science journalism vs. political commentary)
- Time-varying trust scores that account for source reputation changes
- Lightweight cross-referencing when conflicting evidence appears

5.6.5 Handling Contested Claims

The current system struggles with genuinely contested claims where reasonable sources disagree. Future approaches might:

- Explicitly detect and flag claims as “contested” rather than attempting definitive verdicts
- Implement debate-style verification using multiple LLM agents with different perspectives
- Incorporate user feedback to refine handling of ambiguous cases
- Develop specialized pipelines for political versus scientific claims

5.6.6 Production Deployment

Transitioning from research prototype to production would involve:

- Cloud deployment (e.g., AWS with auto-scaling) for handling concurrent users
- User authentication and rate limiting to prevent abuse
- Caching for transcripts and reliability assessments to reduce latency and costs
- Monitoring and alerting infrastructure for operational health
- Browser extension integration for seamless YouTube verification

5.6.7 Extended Evaluation

The current evaluation (30 videos across three domains) should be expanded to:

- A 100+ video corpus with broader topics and video formats
- Multiple annotators to establish inter-annotator agreement baselines

- A richer parameter sweep (higher claim/query limits, more search results)
- User studies evaluating perceived usefulness and trust calibration
- Longitudinal evaluation tracking performance as LLMs and web content evolve

5.6.8 Multimodal Verification

The current system operates only on transcript text, ignoring visual content. Future work could incorporate:

- Image and video frame analysis for visual claim verification
- Chart and graph interpretation for data-driven claims
- Speaker identification and credibility signals
- Deepfake detection for authenticity verification

5.7 Concluding Remarks

This thesis has demonstrated that automated fact-checking of YouTube videos is not only feasible but can achieve practically useful performance levels. Factible represents a step toward democratizing access to fact-checking capabilities, providing users with tools to critically evaluate video content that might otherwise go unexamined.

The system’s strengths—high precision claim extraction, robust evidence retrieval, transparent verdict presentation, and practical efficiency—make it suitable for preliminary fact-checking of YouTube content. Its limitations—particularly with contested political claims—highlight areas where human judgment remains essential and automated systems should complement rather than replace expert fact-checkers.

As misinformation continues to evolve in volume and sophistication, automated fact-checking systems will become increasingly important. The modular, open-source nature of Factible provides a foundation for continued research and development in this critical domain. By making both the system and its evaluation transparent, this work aims to advance the broader goal of building more informed and resilient information ecosystems.

The complete source code, evaluation dataset, and documentation are publicly available at <https://github.com/begoechavarren/factible>, enabling researchers and practitioners to build upon this work.

6 Glossary

This section defines the key terms and acronyms used throughout this thesis.

6.1 Terms

Automated Fact-Checking The use of computational methods to verify the accuracy of factual claims, typically involving claim detection, evidence retrieval, and verdict generation.

Chain-of-Thought (CoT) A prompting technique where a language model is instructed to explain its reasoning step by step before providing a final answer, improving accuracy on complex tasks.

Check-Worthiness A property of claims indicating they are both verifiable (can be confirmed or refuted with evidence) and important enough to warrant fact-checking effort.

Claim Detection The task of identifying factual statements within text that can be verified against external evidence, also known as claim extraction.

Evidence Retrieval The process of finding documents, passages, or data from external sources that support or refute a given claim.

Fact-Checking The process of verifying the accuracy of claims by examining available evidence, traditionally performed by journalists and specialized organizations.

Hallucination A phenomenon where language models generate plausible-sounding but factually incorrect or fabricated information.

Information Retrieval (IR) The field concerned with finding relevant documents or passages from large collections in response to a query.

LLM-as-Judge An evaluation paradigm where a language model assesses the quality of outputs from another model, providing scores or judgments on dimensions like coherence, factuality, or relevance.

Misinformation False or inaccurate information, regardless of intent, that spreads through media and digital platforms.

Multi-Agent System A computational system composed of multiple autonomous agents that interact and collaborate to accomplish tasks, often with specialized roles.

Prompt Engineering The practice of designing and optimizing input prompts to elicit desired behaviors from language models.

Retrieval-Augmented Generation (RAG) An approach that combines language model generation with external information retrieval to ground outputs in retrieved evidence.

Stance Detection The task of classifying the relationship between a piece of evidence and a claim, typically as supporting, refuting, or neutral.

Structured Outputs A pattern in LLM applications where model outputs are constrained to follow a predefined schema (e.g., JSON with specific fields), enabling reliable parsing and validation.

Transcript A text representation of spoken content from a video, including timing information for each segment.

Verdict The conclusion of a fact-checking analysis, typically categorizing a claim as supported (true), refuted (false), mixed, or uncertain.

Zero-Shot Prompting A technique where a language model performs a task based only on the task description, without being provided examples in the prompt.

6.2 Acronyms

API Application Programming Interface

ASGI Asynchronous Server Gateway Interface

CoT Chain-of-Thought

CORS Cross-Origin Resource Sharing

F1 F1 Score (harmonic mean of precision and recall)

FEVER Fact Extraction and VERification (benchmark dataset)

GPT Generative Pre-trained Transformer

HTML HyperText Markup Language

HTTP HyperText Transfer Protocol

IR Information Retrieval

JSON JavaScript Object Notation

LLM Large Language Model

MAE Mean Absolute Error

MAP Mean Average Precision

MBFC Media Bias/Fact Check

ML Machine Learning

MVP Minimum Viable Product

NLP Natural Language Processing

RAG Retrieval-Augmented Generation

REST Representational State Transfer

SDG Sustainable Development Goal

SSE Server-Sent Events

TLD Top-Level Domain

UI User Interface

URL Uniform Resource Locator

WHOIS Protocol for querying domain registration information

YAML YAML Ain't Markup Language (data serialization format)

Bibliography

- [1] Rami Aly, Zhijiang Guo, Michael Schlichtkrull, James Thorne, Andreas Vlachos, Christos Christodoulopoulos, Oana Cocarascu, and Arpit Mittal. FEVEROUS: Fact extraction and verification over unstructured and structured information. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, volume 1, 2021.
- [2] Yochai Benkler, Robert Faris, and Hal Roberts. *Network Propaganda: Manipulation, Disinformation, and Radicalization in American Politics*. Oxford University Press, New York, 2018.
- [3] Wei Chen, Ling Hu, Ming Zhang, and Rui Zhao. LoCal: Logical and causal fact-checking with llm-based multi-agents, 2024. OpenReview preprint.
- [4] Alice Chern, Luis Prieto, and Riya Gupta. A tool-enabled framework for fact-checking language model outputs, 2023. Preprint manuscript.
- [5] Wendy Cheung and Victor Lam. Augmenting llm fact-checking with web retrieval, 2023. Technical report.
- [6] Luyu Gao, Zhuyun Dai, Panupong Pasupat, Anthony Chen, Arun Tejasvi Chaganty, Yicheng Fan, Vincent Y Zhao, Ni Lao, Hongrae Lee, and Kelvin Guu. RARR: Researching and revising what language models say, using language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 16477–16508, 2023.
- [7] Pepa Gencheva, Preslav Nakov, Georgi Karadzhov, Alberto Barrón-Cedeño, and Lluís Màrquez. ClaimRank: Detecting check-worthy claims in political debates. In *Proceedings of the International Conference on Recent Advances in Natural Language Processing*, pages 543–552, 2017.
- [8] Google for Developers. Google fact check tools api. <https://developers.google.com/fact-check/tools/api>, 2024.

- [9] Naeemul Hassan, Bill Adair, James T. Hamilton, Chengkai Li, Mark Tremayne, Jun Yang, and Cong Yu. Detecting check-worthy factual claims in presidential debates. In *Proceedings of the 24th ACM International Conference on Information and Knowledge Management (CIKM)*, pages 1835–1838, 2015.
- [10] Naeemul Hassan, Fatma Arslan, Chengkai Li, and Mark Tremayne. ClaimBuster: The first-ever end-to-end fact-checking system. *Proceedings of the VLDB Endowment*, 10(12):1945–1948, 2017.
- [11] Naeemul Hassan, Gensheng Zhang, Fatma Arslan, Josue Caber, Damian Muthukrishnan, Baoyu Shu, Junghoo Kim, Chengkai Li, and Mark Tremayne. Toward automated fact-checking: Detecting check-worthy factual claims by claimbuster. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1803–1812. ACM, 2017.
- [12] Alan R. Hevner, Salvatore T. March, Jinsoo Park, and Sudha Ram. Design science in information systems research. *MIS Quarterly*, 28(1):75–105, 2004.
- [13] ICWSM Workshop Committee. Proceedings of the icwsm 2025 workshop on agentic fact-checking, 2025. Workshop proceedings.
- [14] Sonu Lakara, Karan Iyer, and Priya Subramanian. MAD-Sherlock: Multi-agent debate for fact verification, 2025. Preprint.
- [15] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pages 9459–9474, 2020.
- [16] Chengkai Li, Naeemul Hassan, Fatma Arslan, Mark Tremayne, and Gensheng Zhang. A platform for live and on-demand monitoring of public discourse. *Proceedings of the VLDB Endowment*, 10(12):1945–1948, 2017.
- [17] Jiawei Li, Han Wu, and Ming Zhou. Automated fact-checking with llm-based claim detection, 2023. Preprint.
- [18] Zheng Lin, Maya Patel, and Alicia Roberts. Fact-Audit: Requirements for trustworthy automated fact-checking systems, 2025. White paper.
- [19] Liang Ma, Shiyu Hu, Wei Zhang, Hang Sun, and Yan Chen. Guided and knowledgeable multi-agent debate for fact verification. *Expert Systems with Applications*, 238:121857, 2025.

- [20] Potsawee Manakul, Adian Liusie, and Mark JF Gales. SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9004–9017, 2023.
- [21] Media Bias/Fact Check. Methodology - media bias/fact check, 2024. URL <https://mediabiasfactcheck.com/methodology/>. Accessed: December 2025.
- [22] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, et al. Webgpt: Browser-assisted question answering with human feedback. *arXiv preprint arXiv:2112.09332*, 2022.
- [23] Angela Ni and Samuel Carter. Structured prompting for consistent claim identification, 2024. Preprint.
- [24] Angela Ni and Samuel Carter. Verifiable claim identification with large language models, 2024. Technical report.
- [25] Briony J. Oates. *Researching Information Systems and Computing*. SAGE Publications Ltd., London, 2006.
- [26] OpenAI. GPT-4 technical report. *arXiv preprint arXiv:2303.08774*, 2024.
- [27] Sebastian Raschka. Llm evaluation: Four practical approaches. <https://sebastianraschka.com/blog/2025/llm-evaluation-4-approaches.html>, 2025.
- [28] Sebastian Ruder. The evolving landscape of LLM evaluation. <https://www.ruder.io/the-evolving-landscape-of-llm-evaluation/>, 2025.
- [29] Marcin Sawiński, Michał Hyben, and Tomasz Wesołowski. Assessing large language models for claim detection tasks. In *Proceedings of the 7th Workshop on Fact Extraction and VERification*, pages 210–221, 2024.
- [30] Timo Schick, Daniel Dwivedi-Yu, Roberta Raileanu, et al. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761*, 2023.
- [31] Skywork AI. How to avoid hallucinations: Editorial fact-check workflow for ai writing. <https://skywork.ai/blog/how-to-avoid-hallucinations-ai-writing-fact-check-guide/>, 2024.
- [32] James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. FEVER: A large-scale dataset for fact extraction and verification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, pages 809–819. Association for Computational Linguistics, 2018.

- [33] Rui Tian, Ming Xie, and Hao Wang. Web-retrieval agents for misinformation detection, 2024. Preprint.
- [34] David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. Fact or fiction: Verifying scientific claims. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7534–7550. Association for Computational Linguistics, 2020.
- [35] Claire Wardle and Hossein Derakhshan. Information disorder: Toward an interdisciplinary framework for research and policy making. Report DGI(2017)09, Council of Europe, 2017.
- [36] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Shaokun Zhang, Erkang Zhu, Beibin Li, Li Jiang, Xiaoyun Zhang, and Chi Wang. Autogen: Enabling next-gen llm applications via multi-agent conversation. *arXiv preprint arXiv:2308.08155*, 2023.
- [37] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *Proceedings of the 40th International Conference on Machine Learning*, pages 1–23, 2023.
- [38] Xuan Zhang, Wei Wei, Yuxiao Wen, Yang Shi, and Bowen Zou. FactAgent: Agentic fact-checking via large language models. *arXiv preprint arXiv:2506.17878*, 2025.
- [39] Yue Zhang, Yafu Li, Leyang Cui, Deng Cai, Lemao Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Yulong Chen, et al. Siren’s song in the ai ocean: A survey on hallucination in large language models. *arXiv preprint arXiv:2309.01219*, 2023.

A Ground Truth Annotation Dataset

This appendix provides details on the ground truth annotation dataset used for system evaluation. The complete dataset comprises 30 YouTube videos with 503 annotated claims across three thematic categories.

A.1 Video Corpus Summary

Table 18 presents the evaluation video corpus organized by category and content type.

Table 18: Evaluation video corpus by category and content type

Video Title	Category	Content Type	Claims
Fossil Fuels: The Greenest Energy	Climate	Misinformation	18
The Great Texas Freeze of 2021	Climate	Misinformation	15
Climate Change: What Do Scientists Say?	Climate	Misinformation	21
Is There Really a Climate Emergency?	Climate	Misinformation	19
Proof: Worldwide Massive Flooding is All Manmade	Climate	Misinformation	25
Causes and Effects of Climate Change (NatGeo)	Climate	Factual	12
Why Does Climate Change Matter	Climate	Factual	9
Extreme Weather	Climate	Factual	14
The Life Cycle of a Plastic Bottle	Climate	Factual	11
What are Greenhouse Gases?	Climate	Factual	13
Fluoridated Water Lowers IQ (Harvard Study)	Health	Misinformation	16
Living with HIV: How Women Were Infected	Health	Misinformation	18
Garlic Cancer	Health	Misinformation	14
3 Detox Juices	Health	Misinformation	12
The Magical 3 Day Juice Fast	Health	Misinformation	15
What Happens When You Exercise Regularly	Health	Factual	17
The Brain-Changing Benefits of Exercise	Health	Factual	15
Immunology Wars: The Battle with HIV	Health	Factual	11
Juice vs. Whole Fruit: Which is Healthier?	Health	Factual	13
What Lack of Sleep Does to the Teenage Brain	Health	Factual	14
Egg Price Warning Comes True	Politics	Misinformation	19
Proof of Election Fraud in 2020	Politics	Misinformation	35
FBI Orchestrated Jan 6th	Politics	Misinformation	22

Continued on next page

Video Title	Category	Content Type	Claims
There is No Gender Wage Gap	Politics	Misinformation	17
A Nation of Immigrants	Politics	Misinformation	16
National Trust Sues Trump Admin	Politics	Factual	20
What Is Democracy (BBC)	Politics	Factual	9
What is Inflation?	Politics	Factual	10
UK Election Results Explained	Politics	Factual	18
Trump’s Historic National Emergency	Politics	Factual	21

A.2 Annotation Schema and Guidelines

Annotations follow ClaimBuster’s check-worthiness principles [11] and store a compact set of structured fields:

- **Claim text:** Verbatim or lightly paraphrased factual statement taken from the transcript; purely opinionated or rhetorical content is excluded.
- **Importance score:** Thesis-impact hierarchy on a 0.0–1.0 scale:
 - 0.85–1.0: Thesis-critical claims whose falsification would collapse the video’s main argument
 - 0.60–0.80: Key evidence claims that significantly support or develop the thesis
 - 0.30–0.55: Contextual claims that provide background or supplementary information
 - <0.30: Peripheral claims of minimal relevance to the main argument
- **Expected verdict:** Anticipated stance label (SUPPORTS, REFUTES, MIXED, UNCLEAR) with a short rationale grounded in established sources. MIXED applies only when credible evidence exists on both sides, while UNCLEAR denotes insufficient or conflicting evidence under the same criteria across all videos.

This schema keeps annotations consistent while capturing the key metadata needed for quantitative evaluation.

A.3 Dataset Statistics by Category

Table 19: Ground truth statistics by category

Category	Videos	Total Claims	Mean/Video	Factual	Misinfo
Climate	10	157	15.7	59	98
Health	10	130	13.0	70	60
Politics	10	216	21.6	78	138
Total	30	503	16.8	207	296

A.4 Importance Score Distribution

Table 20: Distribution of importance scores across ground truth claims

Importance Range	Count	Percentage
High (0.85–1.0)	127	25.2%
Medium-High (0.60–0.84)	198	39.4%
Medium (0.30–0.59)	143	28.4%
Low (<0.30)	35	7.0%

A.5 Expected Verdict Distribution

Table 21: Distribution of expected verdicts across ground truth claims

Expected Verdict	Count	Percentage
SUPPORTS	207	41.2%
REFUTES	189	37.6%
MIXED	68	13.5%
UNCLEAR	39	7.7%

A.6 Data Availability

The complete annotation dataset, including all claims, importance scores, expected verdicts, and video metadata, is available in the project repository at:

<https://github.com/begoechavarren/factible/tree/main/api/experiments/data>

The dataset is provided in YAML format to facilitate both human readability and programmatic access. Each video entry includes:

- Video metadata (YouTube ID, title, category, content type)
- Complete list of ground truth claims with importance scores and expected verdicts
- Tags for filtering and analysis